HEXAWARE

RSA AI Strategy & Use Case Prioritization

Building Momentum & Maturity through a Dual-Path Approach

HEXAWARE

Defining RSA's Path from AI Ambition to Action

RSA partnered with Hexaware to translate AI ambition into executable strategy through an 8-week design sprint — combining **discovery, co-creation, and validation** to define measurable impact, clear governance, and a sustainable path to scale.

Engagement Purpose & Objectives

RSA initiated this work to:

- Identify where **AI delivers measurable business value** aligned to 2025 goals
- **Prioritize and define** the most impactful and feasible AI use cases
- Recommend an **enterprise AI Target Operating Model (TOM)** for scalable, compliant delivery
- Provide **architecture, data, and governance guidance** to enable readiness for adoption

Approach & Outcome

Discovery	Design	Alignment (with stakeholder validation)
Clear view of AI maturity, business priorities, and constraints	Prioritized use case portfolio and TOM recommendation	Leadership alignment and agreed plan for next phase

Validated Through RSA Stakeholder Priorities

Priorities were shaped through RSA-led workshops and reviews, emphasizing **learning over static design**, **scalable and diverse solutions**, **claims and broker analytics as near-term value drivers**, and **data readiness and compliance** as prerequisites — with **industry benchmarks informing, not defining, RSA's path.**

Important Universal Note: This material reflects preliminary analysis and directional targets developed in collaboration with RSA stakeholders. All metrics, financial figures, and outcomes are illustrative and do not constitute validated forecasts. Recommendations are advisory in nature and subject to refinement as data, priorities, and business conditions evolve.

HEXAWARE

How to Use This Strategic Guide — and What's Inside

This guide provides a **structured foundation for RSA's AI journey** — outlining the **strategy, priorities, and enablers** needed to **move from exploration to enterprise adoption**.

Executive Summary: From Promise to Performance	Defines RSA's context, dual-path strategy, and leadership decisions ahead
Dual-Path Strategy in Motion	Introduces the twin approach of Quick Wins ("The Wave") and Strategic Undercurrent ("The Undercurrent")
Target Operating Model (TOM) • Operationalizing the TOM	Defines how AI gets done — including the ways of working, delivery cadence, and RACI needed to operationalize, govern, and scale AI effectively
Use Case Portfolio & Roadmap (RSA validated)	Prioritizes high-value AI use cases and sequences activation milestones that connect early wins to long-term maturity goals
Architecture & Data Enablement	Outlines data, integration, and orchestration enablers that support secure, interoperable, and scalable AI deployment
Appendix	Current State Assessment

Guiding Perspective:

The recommendations in this guide are **high-confidence starting points** for RSA's AI journey — defining **where to focus**, **how to work**, and **how to scale**.

To maximize impact, we recommend a **modern, agile follow-on phase** that advances from **design to disciplined execution** — **testing what works in practice**, **refining through real-world learning**, and **embedding new ways of working** across teams.

Together, these outcomes form the **foundation for next-phase mobilization** — enabling RSA to evolve from **strategy to sustained delivery** with **confidence**, **measurable progress**, and **operational readiness for AI at scale**.

HEXAWARE

Executive Summary:
From Promise to Performance

AS1

Executive Summary: From Promise to Performance

HEXAWARE

RSA's AI Inflection Point: Turning Fragmented Efforts into Focused Advantage

RSA stands at a pivotal moment — moving from isolated AI experiments toward an integrated, value-driven capability that connects strategy, data, and delivery.

Where We Are

Early innovation with pilots and tools in place, but **fragmented governance, data silos, and limited scale.**

Where We're Headed

A unified AI operating model that links **business value to delivery and governance**, embedding AI into everyday processes.

What It Takes

A **dual-path strategy** that pairs quick wins with foundational enablement across process, data, and people.

Slide 5

AS1 This table and overlapping elements should be resized to fit the slide layout

Anna Shymchenko, 2025-10-02T10:55:06.246

AS1

HEXAWARE

Assessing RSA's AI Readiness: Strengths, Gaps, and Priorities for Scale

RSA's AI capabilities show strong intent but remain in the early stages of operational maturity — highlighting the need for disciplined delivery, data enablement, and integrated governance.

Capability Levers

	Incubating	Innovating	Maturing	Optimizing	Transforming
Process	Lack of an organisational AI strategy with a few use cases being experimented, but no view on ROI	Early Innovating — pilots exist, ROI not visible, manual rekeying dominates Start on planning the use of AI in a subset of organisation and POC signed off	AI strategy defined along with Executive sponsorship. ROI expectations are clear	AI strategy execution, aligning with digital, data and technology strategy. ROI metrics defined	AI strategy as a DNA of organisational strategy. AI based operating model and continuous improvement
Data Management	Data Silos across different business units with no clear definition of data requirements for AI	Early Innovating — fragmented datasets, limited lineage, silos persist Clarity on AI driven data requirements with efforts to create a common data store with limited access	Core Data set required for AI solutions is in place. Strategic priority to build data infrastructure for AI strategy	Extensive, upto date and usable data for AI solutions. Strategic systems connect to a common data platform	Data platform fundamental to core business operations. Automated & Self-Service data access
Technology	No investment on an AI technology, however DevOps, RPA and Analytics exists. Lack of clarity on what is needed	Early Innovating — core tools available but not embedded in workflows Ability to provision Cloud or On-premise infrastructure for AI. Model training & deployment still happens manually	AI Deployment architecture and development tools standardised & implemented. Automation of access and resource allocation	AI model deployment increases across BU. Centralised monitoring & audit of AI models for compliance & performance management	State of the art AI infrastructure standardized and efficient. AI used to manage the technology infrastructure itself
People	No defined AI roles & responsibilities. Lack of AI literacy across the business & data teams	Mid-Innovating — AI champions emerging, literacy still low, adoption ~5% of staff Some roles & responsibilities exist but experimentation on right way to organise AI. PoC have started in isolation	Enterprise level AI roles defined. Dedicated CoE created to provide skills. Business leaders communicated on the vision for AI and focus on AI literacy	Clearly defined roles & responsibilities with well-defined KPIs. COE formalised with an organisation wide mandate. Increased AI literacy	High degree of AI literacy with AI first mind set. AI integrated in all roles including executive level. Organisation working model is AI focused
Governance	Employees & leadership being educated about responsible AI. Identification of major risks with priority use cases	Mid-Innovating — framework and council aligned, but not yet operationalized Business, technical & risk team share understanding on responsible AI. Guiding principles of AI adoption beyond legal framework	Guiding principles are converted into specific performance metrics with centralised reporting. Commitment to AI governance is formalised	Formal centralised and auditable process, policy & technology for AI practice. Risk considerations at a model level	Governance goes beyond regulatory compliance. Risk defence and trust are a competitive advantage for applying AI

Stages of Maturity

Slide 6

AS1 This table and overlapping elements should be resized to fit the slide layout
Anna Shymchenko, 2025-10-02T10:55:06.246

Executive Summary: From Promise to Performance

Stable Performance with Headroom to Grow

HEXAWARE

RSA's 2024 results demonstrate stability but reveal a widening efficiency gap compared with top-quartile peers (commercial & specialty focused).

RSA 2024 Performance Baseline

92.8% CR vs ~90% Target

Primary gap : **1.8 - 2.3 CR points** vs commercial focused (w specialty) peers
Strategic Gap: **2.8 CR points** to ~90% Target

Market Context (UK 2024)

Commercial Focused (w Specialty)

AXA XL — **91% CR**

ZURICH UK Commercial — **90.5% CR**

Blended (Commercial & Personal Lines)

Allianz UK incl PL — **95% CR**

AVIVA UK & Ireland incl PL — **94.9% CR**

UK market aggregate: **102.2% CR**
Market softening: 11-20% price reductions

Opportunity Cost

£75-100M Annual value gap (**2.8 CR pts :92.8->90%)**

Q1 2026 AI Opportunity: 0.6 -1.0 CR pts (£16-27 M run-rate)

Strategic Imperative:

Close the performance gap vs commercial peers through pricing discipline, claims excellence, portfolio optimisation, operational efficiency, and ***targeted AI deployment***

DP1

Executive Summary: From Promise to Performance

AI as a Lever to Unlock ~0.6-1.0 CR pt* improvement (≈ £16–27M Annual Run-rate at Adoption)

HEXAWARE

Four RSA validated use cases. Run-Rate realized over **12-18** months (upto **24 months** due to adoption curve + Policy renewal cycles)..

Leverage AI to Close Gap

Target unlock ~0.6-1.0 CR* points improvement starting Q1 2026
Through these four AI intervention clusters

Email/Document Triage for Mailbox

5d → 4hr quote time

- ~0.15-0.2 CR pts improvement target
- 12-14 weeks ; £265K; T-Shirt: (S-M)

Instabase ingestion validation & external data enrichment

Intake → triage → STP
(Intact: 75% pre-fill benchmark)

- ~0.2-0.3 CR basis pts improvement target
- 16-18 weeks; £1.35 -1.75 M; T-Shirt: (L)

Call summarization for claims interactions & sentiment analysis

Claims notes + loss adjuster reports
(Transcript summary)

- ~0.15-0.2 CR basis pts improvement target
- 12-14 weeks; £1 M; T-Shirt: (M)

AI Driven Broker Analytics for Cross Sell

Retention improvement (Regional broker focus alignment)

- ~0.1-0.15 CR basis pts improvement target
- 14-16 weeks; £1.2 M; T-Shirt: (L)

Proof Point:
Intact Canada realized ≈ 4.7 CR pts improvement via 200 + apps (2020-2024) and $150 M+ recurring value through similar use cases (AI+ Pricing+ Claims+ Underwriting)

* CR Targets are indicative, based on industry benchmarks & internal modeling, subject to validation during Q1 2026 execution and 12–24-month adoption curve. Full run-rate benefits realized 2027-2028

Slide 8

DP1 [@Shreyash Tiwari] – we need a different icon here; 2nd and 3rd are the same. Also on slide 40. Thanks.
Donna Pahel, 2025-10-08T00:35:31.266

Executive Summary: From Promise to Performance

Investing to Accelerate Value Realization and AI Capability

HEXAWARE

A phased **£30–40 M** program designed to deliver ROI within 24 months and scale enterprise capability thereafter.

Investment & Returns

Total Investments: £30- 40M
Over 36 months (2026-2028)

Q1 2026 Quick Wins: • Email/Doc Triage for Mailbox: ~£265K {S-M; 12-14 w} • Instabase ingestion validation & external data enrichment:~£1.35-1.75 M {L; 16-18W} • AI Driven Broker Analytics for Cross Sell: ~£1.2M {L; 14-16W) • Call summarization for claims interactions & sentiment analysis: ~£1M {M; 12-14 W)	**~£5M**
2026 Foundation	**~£15M**
2027 Scale	**~£12M**
2028 Optimization	**~£5M**

Industry-Validated Benefits

15-20% Revenue Uplift	**5-15%** Cost Reduction

Incremental gains over 5-year horizon:

Value Horizon	Phase 1 ROI	Adoption
3-5 Years	**2-4X**	**12-24 Months**

Phased Roadmap

Q1 2026: Quick Wins
- Email/Doc Triage for Mailbox (Power UP)
- Instabase ingestion validation & external data enrichment:
- AI Driven Broker Analytics for Cross Sell
- Call summarization for claims interactions & sentiment analysis

Q2-Q4 2026: Foundation
- Scale Q! successes across additional LOB
- Pre-renewal triggers & terrorism question (RSA priority)
- Profin UW Digital assistant+ Premium debiting automation (UKRIS)
- Sanctions & high-risk scanning (RSA compliance)
- Industry opportunity evaluation: Agentic UW, Telematics

2027-2028: Scale & Optimize
- Full STP pipeline across RSA underwriting processes
- Claims AI expansion (voice transcripts, Agentic Claims, STP)
- Industry capabilities: Computer vision, advanced fraud
- Advanced analytics & real-time pricing models

Implementation Reality Check

- Data quality remains critical bottleneck
- Organizational buy-in essential
- Gradual adoption over 12-18 months (upto 24 months)
- Benefits will be incremental, not transformational

NOTE: 58% of insurance CEO's indicate 3-5 years for Gen AI ROI (KPMG 2025); Benefits are incremental (5-15%), Not transformational.

Executive Summary: From Promise to Performance

Learning-First Technical Agenda Anchored in Proven Industry Benchmarks

HEXAWARE

RSA's Q1 use cases span 4 distinct technical integration patterns – building organizational AI literacy and delivery muscle while proving business value across the C&S value chain.

RSA Technical Learning Agenda (4 patterns)

Integration Pattern	Associated Use Case	Description
1) Microsoft Stack	Email/Doc Triage for Mailbox	Power Automate + Outlook + Dynamics integration to accelerate broker response
2) External AI Systems	Instabase ingestion validation & external data enrichment	Instabase AI Hub for intelligent document processing and orchestration
3) Model Training & Development	Call summarization for claims interactions & sentiment analysis	NLP and summarization model development using RSA data to automate report creation
4) GenAI / Agentic Exploration	AI Driven Broker Analytics for Cross Sell	Generative analytics and recommendation models enhance retention and cross-sell

Strategic Rationale

- Build **delivery muscle**: 4 different technical approaches
- Test **Governance**: DPIA, MLOps, explainability across varied risk
- Prove **AI-to-value:** Expense ratio, Loss ratio, revenue, indemnity
- Learn **integration patterns** that scale to future use cases

Industry Benchmark Context*

Positioning Relative to Leading Carriers

Intact Financial
(Parent Company Benchmark)

- 200+ apps
- $150M+ benefits
- 4.7 CR pts (6 years)

UK Competitor: Aviva

- 80+ AI models
- £60M annual savings

US Leaders: AIG, Progressive

- AIG: Agentic UW (1mo → 1d)
- Progressive: Telematics pioneer

Strategic Opportunity Pipeline (2027+)

- **Agentic Underwriting**: AIG-style transformation (evaluate Q2 2025)
- **Telematics Platform:** Progressive benchmark (commercial lines evaluation)
- **Computer Vision:** Tractable partnership opportunity (claims automation)
- **Advanced Fraud:** Multi-modal detection (Allianz/Zurich validated)

Decision Framework:

Phase 1 (Q1 2026)
Execute RSA stakeholder-validated quick wins

>>

Phase 2 (Q2-Q4 2026)
Evaluate industry opportunities based on Phase I learnings (incl Agentic)

>>

Phase 3 (2027+)
Selectively adopt validated industry capabilities for scale (Advanced)

*The AI strategy blends hands-on experimentation with globally validated models from Intact and other leading carriers to accelerate capability, credibility, and competitive parity.

HEXAWARE

Dual-Path Strategy in Motion

Bringing the Dual Path to life — using prioritized use cases to deliver visible wins, build AI delivery discipline, and advance RSA's enterprise learning agenda.

Dual-Path Strategy in Motion

The Dual-Path Approach: Making Waves While Building the Undercurrent

RSA's AI strategy accelerates near-term business impact while laying the groundwork for sustained enterprise capability.

Path 1: Quick Wins 'Making Waves'

Generate visible momentum through focused AI pilots that validate value, delivery discipline, and technical readiness.

- Execute **targeted use cases** to prove value (e.g., Claims Summarisation, Broker Analytics).
- Capture **lessons** on **governance, delivery**, and **architecture** to inform future scaling.
- Build stakeholder confidence and **delivery rhythm** across teams.
- **~5M** investment, **0.6-1.0 CR pts** target

The dual-path approach ensures **momentum and maturity evolve in parallel** — delivering measurable quick-win results while reinforcing RSA's data, governance, and operating model foundations for scale.

Path 2: Strategic Enablement 'The Undercurrent'

Build the architecture, governance, and enterprise capability that sustain and scale AI value.

- Strengthen **data and model management foundations** to support enterprise-scale AI
- Formalize **governance frameworks** and **CoE operating structures**.
- Enable **long-term**, **scalable**, and **compliant AI capability** across business functions.

Dual-Path Strategy in Motion

Path 1: Quick Wins — Proving Value and Building Delivery Muscle

HEXAWARE

Quick Wins focus on turning AI ambition into action — using targeted pilots to prove value, strengthen delivery muscle, and establish the operating foundations for scale.

Path 1 Focus (Q4 2025 – Q1 2026)

Primary Objectives:

- **Demonstrate AI-to-value linkage** across top-priority use cases
- **Build delivery discipline** through rapid, measurable AI deployments
- **Test governance and operating model MVP** under real delivery conditions
- **Strengthen cross-functional collaboration** across IT, business, and data teams

Secondary Objectives:

- Capture **lessons to inform enterprise architecture** and future scaling
- Establish early **success metrics for credibility and internal communication**

Quick Win Use Cases

Email Triage for Mailbox (Microsoft Stack)

- 5d → 4hr quote time; UW productivity
- Automate intake and routing using Power Automate + Outlook + Dynamics; Human-in-the-loop (**12-14 w; £ 265K)**

Instabase ingestion validation & external data enrichment

- 80% + auto validation target; 100% sanctions checks
- Segment 1: Post-Instabase validation (Q1 26); Segment 2: Enrichment (Q2-Q3 26)
- Current: Commercial Combined only; expand to all Lob's (**16-18 w; (two phases);£ 1.35M-1.75 M)**

Call summarization for claims interactions & sentiment analysis (NLP)

- Real time transcription + AI summary of claims notes and loss adjuster transcripts+ compliance prompts
- Improve adjuster efficiency and claim-cycle visibility **(12-14 w; £ 1.0 M)**

AI-Driven Broker Analytics for Cross Sell (CRM / GenAI)

- **19+**broker portals; propensity scoring; CEO/ Regional escalation framework
- Retention improvement (10-12%)+ Cross sell acceleration
- Increase alignment between regional broker strategy and performance **(14-16 w; £ 1.2 M).**

Expected Outcomes

By end of Q1 2026:

- **£16–27 M run-rate benefit (0.6-1.0 CR pts)**
- **Realized over 12-18 months (upto 24 months); Q1 2026 impact £ 5-10M; Full run-rate 2027-28**
- **Validated AI-to-value patterns (4)** ready for scale
- **Governance tested;** DPIA, MLOps, explainability across varied risk profiles
- **Improved stakeholder confidence** through visible early wins + reusable integration patterns
- **Refined TOM foundations** for data, process, and governance readiness

***Note:** Impact estimates are directional, informed by Intact and peer benchmarks. Realized benefits will depend on RSA's data, adoption, and delivery maturity.*

Dual-Path Strategy in Motion

Path 2: Strategic Enablement — Building the Undercurrent

The Undercurrent" moves RSA from frameworks to flow — embedding data, governance, and operating-model discipline into business-as-usual so that AI becomes systematic, not episodic.

Foundational Priorities (2026-2028)

Data & Architecture

- Strengthen **data governance** and **quality frameworks** aligned with Intact enterprise principles
- Implement **EDP integration** and standardized data models for cross-function visibility
- Create a **central feature store** and metadata registry to enable model re-use and consistency

AI Governance & Operating Model

- Operationalize RSA's **AI governance framework** into daily workflows, decision gates, and delivery routines
- Formalize a **Center of Excellence (CoE)** with federated domain teams for sustainable delivery
- Define and implement **model lifecycle management** (monitoring, retraining, bias testing, audit readiness).
- Accelerate **DPIA completion:** Critical completion gate for Q1 2026 deployment.

Enterprise Capability & Change

- Build **cross-functional enablement** through AI literacy and process redesign
- Implement **value-tracking and measurement frameworks** to sustain ROI visibility
- Create an **innovation pipeline** for continuous improvement and talent development.

+ Key Enablers + Initiatives

By end of 2028:

Data Foundation & Architecture

- Align RSA data governance to Intact enterprise standards
- Define common data models, feature stores, and metadata registries

Governance & Operating Model

- Translate RSA's AI governance framework into day-to-day delivery controls
- Stand up AI Center of Excellence (CoE) to manage pipeline and standards

Capability Enablement

- Deploy change and training programs to scale AI skills and delivery discipline
- Establish innovation and model retraining pipelines for continuous improvement.

+ Expected Outcomes

By end of 2028:

- **Governance and data frameworks** fully operationalized, embedded in day-to-day delivery
- **Scalable AI** capability embedded across underwriting, claims, and broker enablement
- **Compliance and transparency** fully aligned to FCA and EU AI Act requirements
- RSA achieves **competitive parity with leading UK insurers (Aviva, Zurich, AXA) on AI Maturity**
- **2.0-2.5 CR pts** cumulative improvement via integrated transformation

Note: *Strategic enablement priorities build on RSA's existing governance foundation and align with Intact enterprise data principles. Outcomes and timing will depend on data readiness, adoption velocity, and cross-functional alignment through 2026–2028.*

HEXAWARE

Target Operating Model (TOM)

TOM

Target Operating Model (TOM) — How AI Gets Done at RSA

A disciplined operating model defining *how AI gets delivered, governed, and scaled* —embedding agility, compliance, and collaboration across RSA.

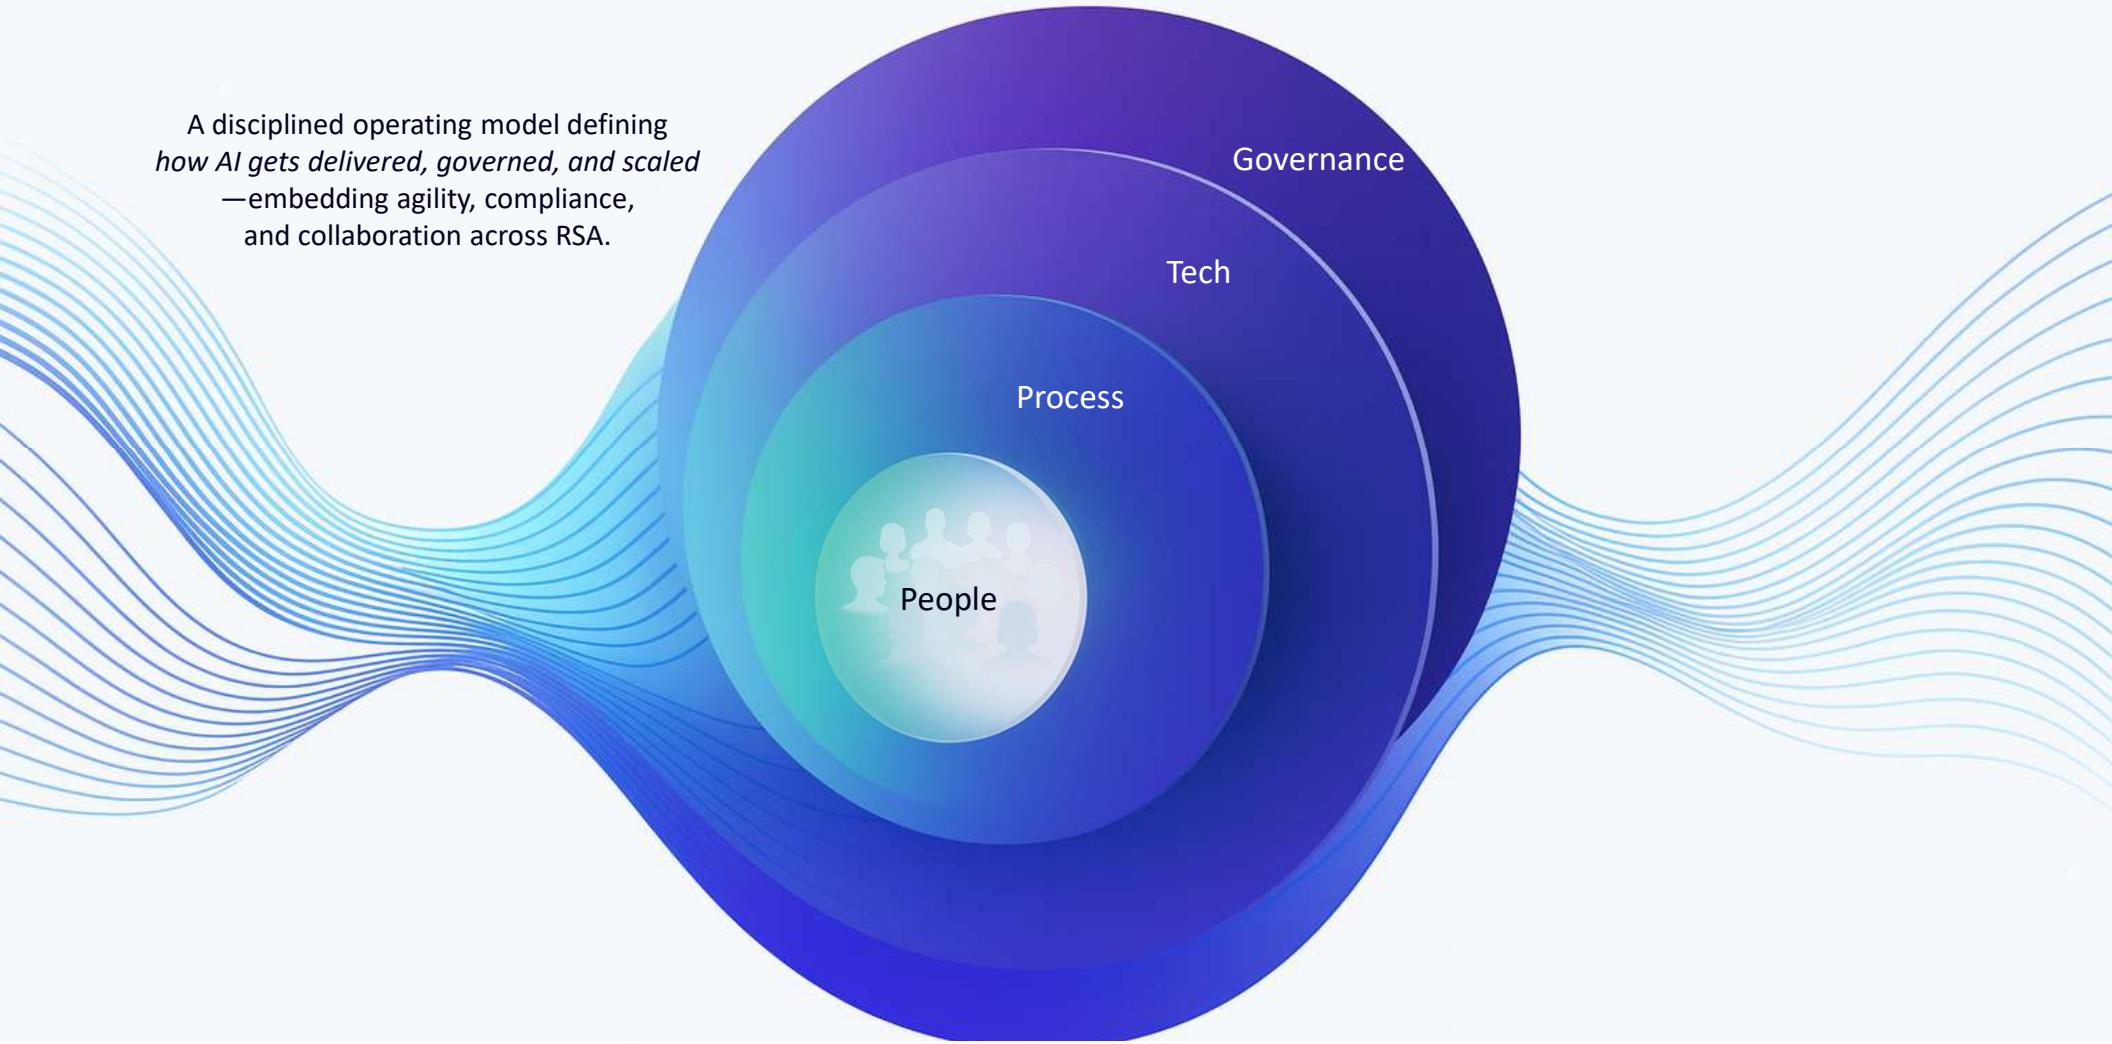

RSA AI Adoption Principles – Aligned with Dual-Path Strategy

HEXAWARE

These principles translate the dual-path strategy into actionable disciplines — embedding learning, governance, and enterprise integration into every AI initiative.

Foundation

Learning Agenda + Governance-by-Design + Enterprise Integration

Business Value + Learning Approach	Adopt Agile and Iterative Approaches	Governance by Design (Not Afterthought)	Data Foundation + Enterprise Platform Integration	CI/CD + MLOps for Production Readines	Documentation + Knowledge Transfer
Align AI initiatives with RSA business objectives (risk-free migration, operational efficiency, broker/customer experience) while prioritizing organizational learning over immediate financial returns.	Break down AI delivery into manageable sprints/phases using RSA's established COE model with cross-functional pods empowered with P&L accountability.	Embed governance, ethics, and compliance from inception through established Ethical SteerCo and Working Group frameworks, with DPIA as regulatory gate.	Connect AI initiatives to Enterprise Data Model (EDM) and strategic platforms (EDP, Power UP, Guidewire) to avoid tactical "Unicorn" solutions and technical debt.	Implement continuous integration and deployment pipelines with model operations (MLOps) to automate build, testing, deployment, and monitoring at scale.	Deliver system documentation with every AI workstream to enable organizational learning, support scale, and build institutional capability.

Embedded governance through DPIA compliance, Ethical SteerCo oversight, and regulatory readiness (FCA/GDPR/EU AI Act) | Built on established COE model with cross-functional pods

From Idea to Impact — The AI Delivery Workflow

The operating model translates into a validated, end-to-end workflow — guiding every AI initiative from ideation through scaled deployment under a consistent governance and delivery rhythm.

Ideate	Evaluate	Prioritise	Solution Design	Build & Test	Deploy	Scale & Optimise
• Inside-out (business pain) • Outside-in (industry best practice) • Intact playbook reference • Strategic program alignment	• 10-lever scoring (impact + effort) • Ethical review • Technical feasibility • Data readiness assessment	• Quadrant placement • Strategic alignment • Resource allocation • Portfolio balancing	• Solution architecture • User cases/stories • Integration patterns • Data model design	• Agile sprints (2 weeks) • Model development • SIT/UAT execution • DPIA completion	• Cutover execution • Production deployment • Hypercare support • Performance monitoring	• Enterprise rollout • Continuous improvement • Model retraining • Lessons learned

Cross Functional POD

Business Product Owner
Prioritizes value, validates workflows, bridges business and IT

CIO Team (COE)
Delivers AI platforms, defines architecture, enables reuse

Domain SMEs:
Provide rules, validate logic, support UAT

QA/ Test Lead
Plans tests, ensures quality, coordinates performance/VAPT

Scrum Master
Facilitates sprints, removes blockers, ensures agile delivery

Data Scientists:
Build, validate, and deploy ML models

Enterprise Architecture
Reviews design, ensures roadmap alignment, prevents tech debt

Change Management
Drives adoption, trains users, manages workflow transitions

Slide 18

AM1 the same comment about labels/terms consistency: I wuld make it nouns everywhere (ideate -> ideation, deploy-deployment etc)

Anton Makarevich, 2025-10-06T09:04:08.098

Governance with Discipline to Drive AI Success

HEXAWARE

A multi-layered governance model that embeds compliance and ethics into delivery — ensuring AI moves from intent to measurable, risk-managed performance.

Core Governance Principle

Governance that is *"seamless for business to adopt"* — enabling accountability without friction.
The right governance, at the right depth, ensuring compliance, risk control, and delivery speed in equal measure.

Ethical Steering Committee Chair: UK&I Chief UW Officer (SMF23)

- Approve guiding principles for ethical AI use
- Review/challenge AI initiatives for customer outcomes
- Final approval authority for high-impact AI
- Six-monthly reporting cadence
- FCA/GDPR/EU AI Act compliance oversight

Ethical Working Group Chair: Chief Technical UW Officer)

- Review AI initiatives at inception
- Mandatory DPIA sign-off before production
- Business rules and ethical impact assessment
- Monthly meeting cadence
- Escalate high-risk cases to SteerCo
- Current Gap: Only 2 of 48 initiatives have completed DPIA

Program Steering Committee Program-Level Governance

- Strategic programs own AI use cases (CL Tech & Workflow, CL Pass, Guidewire, Symphony)
- Monthly steering committee meetings
- Quarterly gate reviews for major milestones
- Budget and resource allocation
- Cross-program dependency management

Change Control Board Technical & Operational Governance

- Evaluate, approve, and track changes to AI systems
- Model version control and retraining governance
- Integration testing and deployment approval
- Post-deployment performance monitoring
- Incident response and remediation oversight

Ideate	Evaluate	Prioritise	Solution Design	Build & Test	Deploy	Scale & Optimise
Business Product Owner + CIO	**Ethical Working Group**	**Program Steering Committee**	**Change Control Board + Working Group**		**Ethical Working Group**	**Ethical SteerCo** (if high-impact)
• Use case submission • Alignment to strategic priorities • Preliminary feasibility assessment	• DPIA initiation • Ethical impact assessment • 10-lever scoring validation • Data privacy review	• Portfolio placement decision • Resource allocation approval • Budget confirmation • Dependency mapping	• Architecture review • DPIA completion sign-off • Model validation and testing • Integration approval		• Production readiness review • Final DPIA sign-off, • Deployment authorization Monitoring plan approval	• Enterprise review • Compliance validation Performance reporting Bi-Annual SteerCo updates

Key Risks & Mitigations for AI Initiatives **(1/2)**

HEXAWARE

RSA's AI Maturity Assessment surfaced critical risks across compliance, model accuracy, data integrity, and SME bandwidth — each addressed through targeted, embedded mitigations.

Risk Category	Mitigation Measures
Regulatory Compliance (FCA, GDPR, EU AI Act)	• Establish framework to identify regulatory change and manage compliance proactively • Mandatory DPIA completion for all AI initiatives before production • Ethical SteerCo and Working Group review aligned to FCA/GDPR requirements • Quarterly compliance assessments and regulatory horizon scanning • Distinct UK/EU compliance requirements (not Intact Canada framework)
Model Risk & Accuracy (Drift, bias, explainability)	• Implement pre-implementation and ongoing model testing frameworks • Increased supervision with "human in the loop" for high-stakes decisions • Use training models to recognize normal/abnormal patterns and generate alerting • Address current MLOps maturity gaps (model retraining, drift monitoring) • Establish model performance tracking and validation cadence
Data Quality & Privacy (PII, accuracy, completeness)	• Implement Dynamic Data Obfuscation & Redaction controls to reduce data loss risk • Implement feedback loops and remediation to manage data accuracy and bias • Connect AI initiatives to Enterprise Data Model (EDM) and EDP roadmap • Address fragmented data landscape (8 policy admin systems, 2 CRMs) • Enforce data privacy, access controls, and retention policies through IT/Data Governance
SME Availability (Domain expert bandwidth)	• Publish stakeholder engagement plan and prioritize use cases based on SME availability • Leverage AI-powered accelerators to minimize SME time requirements • Secure timely access to SMEs for validations and critical decisions • Two-in-a-box model: BA and tester work together to validate insurance-specific logic • Focus on insurance validation points: capital evolution, risk rating, premium calculation

AM1

Key Risks & Mitigations for AI Initiatives **(2/2)**

HEXAWARE

Operational, adoption, and ROI risks are managed through technical co-existence strategies, proactive change enablement, and clear cost-to-value tracking.

Risk Category	Mitigation Measures
Technical Complexity & Debt (Integration, legacy systems)	• Build co-existence strategies for legacy and modernized systems during transition • Adopt loosely coupled architecture design to minimize integration impact • Avoid tactical "Unicorn" solutions - align to Architecture Blueprint strategic platforms • Don't increase reliance on systems to be decommissioned (legacy by 2028 target) • Embed AI in DevOps with CI/CD pipelines and automated testing
Adoption & Change Management (User resistance, training)	• Leverage OCM function to ensure organizational preparedness and plan necessary skilling • Prepare business users with role-based training (agents, brokers, adjusters, underwriters) • Drive adoption through workshops, pilot rollouts, and structured feedback loops • Broadcast and celebrate milestones to boost morale and engagement • Front-line adoption by distributed users is critical - get early feedback to reduce post-launch changes
Cost Overruns & ROI (LLM costs, benefit realization)	• Onboard models into FinOps process; review and adapt for Gen AI cost models • Identify charging model for AI tools and agree tooling as part of AI roadmap • Learn from initial use cases to improve cost estimation capability (review PoCs for accuracy) • Right-size model selection and licensing to minimize duplication of capabilities • Track actual ROI against projections and feed back to Value Framework
Talent & Skills Gap (AI expertise retention)	• Define sourcing strategy for AI talent (build vs. buy vs. partner) • Agree standards for documentation and knowledge transfer to build institutional capability • Invest in automation to reduce dependency on scarce skills • Consider AI upskilling and awareness programs across the organization • Only host models in-house where necessary to manage resource/skill requirements

Slide 21

AM1 should we mention security of AI solutions?
Anton Makarevich, 2025-10-03T18:13:56.429

AM1

TOM

Best Practices for AI Delivery at Scale — Embedding Excellence in Every Pod

HEXAWARE

Proven practices from RSA's delivery muscle — blending agile cadence, enterprise integration, and governance-by-design to ensure repeatable success.

Learning-Led Approach
- Early-on prioritise learning value over immediate financial benefit
- Challenge business thinking and expand AI horizons
- Build cross-functional design and delivery muscles
- Diverse technical learning across integration patterns
- Test governance and operating model MVP
- Validate AI-to-value linkage for stakeholder credibility

Leverage COE Model
- CIO organization pioneers technology introduction
- Pilot with business partners (not IT working in isolation)
- Deliver common AI capabilities at scale
- Hub-and-spoke for Data & Analytics
- Platform teams provide reusable integration patterns
- Cross-functional pods with P&L accountability

Governance by Design
- Embed governance from inception (not afterthought)
- DPIA as mandatory regulatory gate
- FCA/GDPR/EU AI Act compliance from Day 1
- Governance that is seamless for business to adopt
- Risk assessment at appropriate level (use case vs. capability)
- Route through Ethical Working Group → SteerCo

Enterprise Integration
- Connect to EDM and EDP roadmap
- Leverage strategic platforms: Power UP, Instabase, Radar Live, Guidewire
- Avoid tactical "Unicorn" solutions and technical debt
- Reduce reliance on systems to be decommissioned
- Move from tactical data marts to enterprise architecture
- Address system fragmentation through consolidation

Agile Delivery Cadence
- Deliver in demonstrable 8-12 week increments
- 2-week Agile sprints with clear success criteria
- Monthly steering committee + quarterly gate reviews
- Continuous sprint demos and stakeholder feedback
- Plan for dress rehearsal before production cutover
- Hypercare support post-deployment

Define Clear Acceptance Criteria
- Establish acceptance and exit criteria for each phase
- Align with business objectives per RACI matrix
- Two-in-a-box validation (BA + tester for insurance logic)
- Focus on insurance-specific validation points
- UAT with front-line users before enterprise rollout
- Track actual outcomes vs. projections

Plan SME Engagement
- Publish SME engagement plan early
- Secure timely access for critical decisions and validations
- Prioritize workstreams based on SME availability
- Leverage AI accelerators to minimize SME time
- Domain SMEs define how system behaves for each LOB
- Early involvement prevents late-stage rework

Invest in Test Environments
- Execute PoC/PoT to validate and establish confidence
- Test automation to ensure comprehensive coverage
- Establish pre-prod that mimics production for NFR validation
- Ensure adequate capacity for parallel workstreams
- Address gaps in automated validation and testing
- Build robust Guidewire AI pipeline (currently not mature)

Comprehensive Documentation
- Business rules, functional specs, solution architecture
- Test strategy, test cases and data, automation scripts
- Model documentation (algorithms, training data, metrics)
- Integration patterns and reusable components catalogued
- KT sessions for organizational learning
- Consider AI upskilling programs across organization

One Commercial Guiding Principle (Applied to AI)

"Speed is critical. One value proposition allows us to simplify processes, mix teams and eliminate 'RSA way, NIG way' conversations and move onto the 'new way'"

Applied to AI: Eliminate siloed approaches, create unified AI delivery model, accelerate through standardization, enable "One AI Way" across RSA.

Slide 22

AM1 this slide is very difficult to read, also design is very different from similar slides
Anton Makarevich, 2025-10-06T09:05:53.319

AM1 0 can we use the same panels as on the next one?
Anton Makarevich, 2025-10-06T09:06:37.836

Clarifying RSA Stakeholder Roles — Part 1: Business, Platform, and Governance

HEXAWARE

Defining who owns what ensures AI delivery aligns with business value, domain logic, and compliance from the start — not as an afterthought.

Business Product Owner
(e.g., Head of Claims, UW, or Policy Admin)

Responsibilities
- Prioritizes features based on business value
- Validates user stories and functional workflows
- Acts as bridge between IT and operations
- Defines success criteria and acceptance criteria
- Final sign-off on UAT and production readiness

Why Critical
Anchors the AI solution in real insurance business processes (e.g., FNOL, premium calculation, endorsements, claims triage). Without PO ownership, AI becomes technology looking for a problem.

Domain SMEs
(Claims Analysts, Underwriters, Actuaries)

Responsibilities
- Provide business rules for rating, reserving, policy issuance, claims triage
- Verification of business rules and workflow processes
- Participate in gap analysis and UAT
- Validate outputs for accuracy and reasonableness

Why Critical
These experts define how the AI system should behave for each LOB. AI models trained without SME validation risk incorrect business logic, poor risk selection, and regulatory non-compliance.

CIO Organization
(Platform Teams) - COE Model

Responsibilities
- Pioneer technology introduction & pilot with business
- Deliver common AI capabilities (Copilot, Instabase)
- Provide technical architecture & integration patterns
- Ensure alignment to Architecture Blueprint and strategic platforms

Why Critical
RSA's established COE model for tech delivery at scale. CIO pioneers then pilots with business partners - not working in isolation. Critical to avoid tactical solutions and technical debt.

Data & Analytics Team
(Hub-and-Spoke Model)

Responsibilities
- EDP integration, data quality frameworks, analytics infrastructure
- Data scientists embedded in business teams
- ML model development, validation, and monitoring
- Connect AI initiatives to Enterprise Data Model (EDM)

Why Critical
Central platform + federated BI/MI ensures data foundation for AI. Hub-and-spoke model embeds data scientists in business context while maintaining enterprise architecture standards.

Enterprise Architecture
(Graham Wilde & team)

Responsibilities
- Solution architecture & integration pattern validation
- Alignment to Architecture Blueprint & strategy
- Avoid tactical "Unicorn" solutions and technical debt
- Ensure AI integrates with modernization roadmap (decommission legacy by 2028)

Why Critical
Prevents AI initiatives from creating technical debt or increasing reliance on systems to be decommissioned. Ensures diverse technical learning across integration patterns.

Ethical SteerCo + Working Group
(Andrew Tongue & AI Council)

Responsibilities
- Approve guiding principles for ethical AI use
- Review/challenge AI initiatives for customer outcomes and regulatory compliance
- Mandatory DPIA sign-off prior production deployment
- FCA/GDPR/EU AI Act compliance embedded from start

Why Critical
Insurance is highly regulated. Governance-by-design ensures compliance, ethical use, and avoids reputational/regulatory risk. DPIA is regulatory gate - non-negotiable.

Clarifying RSA Stakeholder Roles — Part 2: Adoption, Change, and Compliance

HEXAWARE

Frontline adoption, data governance, and compliance are integral to AI success — requiring coordinated ownership across IT, risk, and change functions.

IT & Data Governance Leads

Responsibilities

- Integrate with core system (PAS, CRM, Guidewire, EDP)
- Infrastructure readiness (cloud/on-prem deployment)
- Data privacy, access controls, retention policies
- Security tooling to support AI processes (Dynamic Data Obfuscation, PII controls)

Why Critical

Insurance deals with legacy tech, complex integrations, and regulated data (PII). Governance ensures data security, privacy compliance, and operational resilience.

Change Management & Training

Responsibilities

- Prepare business users for new AI tools and workflows
- Create role-based training (agents, brokers, adjusters, underwriters, claims handlers)
- Drive adoption through workshops, feedback, & pilots
- OCM for cultural transformation

Why Critical

Front-line adoption by distributed users (brokers, agents, adjusters) is critical to AI success. Success depends on thousands of users embracing new ways of working - especially in broker-led channels.

End Users (Agents, Brokers, Claims Adjusters, Underwriters)

Responsibilities

- Participate in design validation, usability testing, and pilot rollouts
- Report pain points & suggest improvements in UAT
- Provide feedback on AI outputs and model accuracy
- Champion adoption within their teams

Why Critical

Primary users - getting their feedback early ensures better adoption and fewer change requests post go-live. Front-line input prevents building AI that doesn't meet real-world needs.

Compliance & Risk Officers

Responsibilities

- Review AI initiatives for regulatory compliance
- Approve use cases, user stories, and functional workflows from compliance lens
- Ensure model explainability and audit trails
- Validate that AI doesn't discriminate

Why Critical

Insurance is highly regulated, and compliance must be built into the AI lifecycle. Prevents regulatory exposure, reputational damage, and ensures consumer protection.

Critical Success Factor

Multi-party collaboration is essential for AI delivery success.

- RSA stakeholders own business context, domain expertise, governance, and change management.
- Delivery partners provide technical execution, but cannot succeed without active RSA engagement across all these roles.

> **"AI is at the top of our list...from overall organization perspective...this requires whole-organization commitment, not just IT project execution." – RSA Stakeholder**

HEXAWARE

AI Delivery Work Breakdown Structure (WBS) — From Ideation to Deployment (1/2)

The full delivery rhythm of AI programs — mapping every phase of the TOM to governance gates, deliverables, and stakeholder actions.

Program Management (Horizontal Activity Across All Phases)

Project Plan & Deliverables Tracking | **Status Updates** (WSR) | **Risk Assessment** | **Stakeholder Management**

Ideate	Evaluate	Prioritise	Solution Design	Build & Test	Deploy	Scale & Optimise
Program Management	**Ethical Working Group DPIA Initiation**	**Program Steering Quadrant Placement**	**Enterprise Architecture Review**	**DPIA Completion & Sign-off**	**Final DPIA + Change Control Board**	**Ethical SteerCo (if high-impact)**
Business pain workshops (inside-out)	10-lever scoring (Impact + Effort)	Quadrant placement decision	Solution architecture design	Agile sprints (2-week cycles)	Cutover plan approval	Enterprise rollout planning
Intact playbook review (outside-in)	DPIA kickoff (mandatory)	Resource allocation approval	User stories & acceptance criteria	Model development & training	Production deployment	Model retraining pipelines
Strategic program alignment check	Ethical impact assessment	Budget confirmation (T-shirt sizing)	Data model design (EDM connection)	SIT execution & defect resolution	Hypercare support (2-4 weeks)	Drift monitoring & alerting
Initial use case submission	Data readiness check	Cross-program dependency mapping	Integration patterns (Instabase/Guidewire)	UAT with domain SMEs (5-10 users)	Performance monitoring setup	Knowledge transfer & lessons learned
Preliminary feasibility assessment	Model risk classification	Pod formation authorization	Insurance business rules definition	DPIA completion & sign-off (mandatory gate)	User training & guides	Continuous improvement

Agile Delivery Cadence & Governance Rhythm

Sprint Demo (2-week cycles)	Daily standup	Monthly Steering Committee	Quarterly Gate Reviews

AI Delivery WBS — Roles, Validations, and Governance Gates (2/2)

A cross-functional pod structure with P&L accountability ensures that every phase of delivery is validated through mandatory governance and insurance-specific checkpoints.

Cross-functional pod structure with P&L accountability

- **Business Product Owner** (Claims/UW/Policy Admin Lead)
- **Scrum Master** (Agile Facilitator)
- **CIO Platform Team**
- **Data Scientists** (hub-and-spoke)
- **Domain SMEs** (UW/Claims/Actuaries)
- **Enterprise Architecture**
- **QA/Test Lead**
- **Change Management**

8-12 Person Team comprising of above roles

Governance Gates & Insurance Validations

Mandatory Gates

- Phase 2: DPIA Initiation (Ethical Working Group)
- Phase 3: Portfolio Placement (Program Steering)
- Phase 5: DPIA Completion (Ethical Working Group)
- Phase 6: Final DPIA + Deployment Auth (Change Control Board)

Insurance-Specific Validations

- Rating algorithms: Premium calculation accuracy
- Underwriting logic: Risk scoring, auto-referral rules
- Claims triage: FNOL automation, loss ratio validation
- Policy issuance: Endorsements, renewals, cancellations
- Capital evolution: Reserve adequacy, combined ratio impact
- Broker experience: Quote speed, submission pre-fill

MLOps & Model Governance

- Model performance baseline established
- Drift detection thresholds configured
- Retraining triggers documented
- Model versioning and lineage tracked
- Explainability framework for audit trail

Transition to Operations

- System handover to RSA BAU teams
- Integration patterns documented for reuse
- Actual ROI vs. projections tracked
- Feedback loop to Value Framework
- Post-implementation review completed

Mandatory governance gates at Phases 2, 3, 5, 6

Operationalizing the TOM

*Next, let's bring our TOM **from concept to practice** — demonstrating how **governance**, **delivery rhythm**, and **cross-functional roles** come together in action through a **real or representative AI use case**.*

***The aim:** to show **how AI gets done at RSA** — not just how it's designed to.*

AM1

Operationalizing the TOM

Stage 1: Ideate — Defining the Right AI Opportunities

HEXAWARE

RSA's ideation process blends inside-out business pain discovery with outside-in market intelligence to ensure AI initiatives target real operational bottlenecks while aligning with industry best practice and regulatory readiness.

Ideate | Evaluate | Prioritise | Solution Design | Build & Test | Deploy | Scale & Optimise

Inside-Out (Business-Driven)

- **Strategic Programs:** One Commercial, Symphony, CL Tech & Workflow, CL Pass, Guidewire rollout, Integrated Platform Pricing
- **Business Pain Points:** Workshops with UW, Claims, Broker teams identifying operational bottlenecks and manual process inefficiencies
- **Regulatory/Compliance Needs:** FCA Consumer Duty requirements, GDPR compliance, EU AI Act readiness, Ethical SteerCo mandates
- **Front-Line Feedback:** Underwriters, Claims Adjusters, Brokers reporting pain from dual-key data entry, manual routing, Excel workarounds
- **Executive Priorities:** ~90% CR target, £2.7B→£5B DPW by 2030, broker value proposition improvement, claims efficiency
- **Current AI Inventory:** 48 initiatives in Risk Register requiring governance alignment and DPIA completion
- **Architecture Pain Points:** 8 policy admin systems, 2 CRMs, fragmented data landscape, legacy decommission by 2028 target

Outside-In (Industry Best Practice)

- **Intact Group Playbook:** 200+ AI apps deployed since 2018, $150M+ recurring benefits, 4.7pts CR outperformance – adapted for UK/EU regulatory context
- **Intact Target Benchmarks:** 75% submission pre-fill by 2026 (Canada achieved 2024), automated quote parsing with LLM inference, data enrichment with risk hazards/firmographics
- **Technology Vendor Capabilities:** Instabase AI Hub roadmap, Power Platform AI features, Guidewire AI integration patterns, Radar Live pricing enhancements
- **Industry Research:** Everest Group prioritization frameworks, Gartner insurance tech trends, FCA thematic reviews on AI usage
- **Competitor Intelligence:** Market scanning for AI-driven speed-to-quote, automated underwriting, digital claims, broker analytics
- **Emerging Technologies:** Agentic AI capabilities, LLM advancements (GPT-4, Claude), MLOps maturity frameworks, enterprise AI governance standards
- **Hexaware Portfolio:** Insurance-specific AI use cases across underwriting, claims, pricing, policy servicing value chain

Slide 28

AM1 we agreed to go with 7 stages everywhere, right?
Anton Makarevich, 2025-10-06T09:07:21.261

Operationalizing the TOM

Stages 2 & 3: Evaluate and Prioritise — Balancing Impact and Effort

Each AI use case is scored through a 10-lever framework that quantifies business value, feasibility, and risk, allowing RSA to prioritise initiatives that deliver the greatest impact with manageable implementation effort.

Ideate | **Evaluate** | **Prioritise** | Solution Design | Build & Test | Deploy | Scale & Optimise

RSA AI Value Framework: 10-Lever Scoring System

Every use case evaluated across 10 dimensions (1-5 scale) to determine **Business Impact Score** and **Implementation Effort Score**

Business Impact (5 Levers - Equally Weighted at 20%)

1. **Revenue Impact:** Direct revenue growth (pricing, new products, policy growth) | 5 = £5M+ annually
2. **Cost Savings:** Operational efficiency, automation, fraud reduction | 5 = £5M+ annually
3. **Risk Reduction:** UW accuracy, compliance, fraud detection, regulatory | 5 = Critical risk elimination
4. **Broker/Partner Experience:** Quote speed, service quality, relationship strength | 5 = Transformative improvement
5. **Strategic Fit:** Alignment to One Commercial, ~90% CR target, modernization | 5 = Core to digital transformation

Implementation Effort (5 Levers - Equally Weighted at 20%)

1. **Data Readiness:** Availability, quality, integration needs | 5 = Data doesn't exist or major infrastructure needed
2. **Technical Complexity:** AI sophistication, system integration, novel algorithms | 5 = Cutting-edge agentic AI
3. **Change Impact:** Organizational adjustment, process redesign | 5 = Enterprise-wide transformation
4. **Model Risk:** Regulatory scrutiny, explainability, customer impact | 5 = High-stakes automated decisions
5. **Adoption Readiness:** User eagerness, training needs, cultural barriers | 5 = Major resistance/change management

Scoring Threshold:

3.0 (configurable) | Quadrant Logic: Impact ≥3.0 & Effort <3.0 = Quick Win | Impact ≥3.0 & Effort ≥3.0 = Strategic Bet | Impact <3.0 & Effort <3.0 = Experimental | Impact <3.0 & Effort ≥3.0 = Watchlist

Strategic Quadrant Placement

QUICK WINS
High Impact | Low Effort
Priority: Immediate deployment (40% portfolio allocation)
Example: Email triage, document extraction, broker insights

STRATEGIC BETS
High Impact | High Effort
Priority: Planned investment (30% portfolio allocation)
Example: E2E Agentic UW, predictive pricing, claims AI pipeline

EXPERIMENTAL
Low Impact | Low Effort
Priority: Learning & innovation (20% portfolio allocation)
Example: ChatGPT search tools, PoC, technology pilots

WATCHLIST
Low Impact | High Effort
Priority: Re-evaluate or defer (10% portfolio allocation)
Requires business case refinement or market condition changes

Governance Gate

- All use cases must pass through Ethical Working Group → Ethical SteerCo review.
- DPIA completion is mandatory before moving to delivery.
- **Current gap: Only 2 of 48 initiatives have completed DPIA.**

Operationalizing the TOM

Stage 4: Deliver — Cross-Functional Roles in Action (Part 1)

HEXAWARE

AI delivery at RSA operates through cross-functional pods that unite business owners, domain SMEs, data teams, and platform leads—ensuring clear accountability and rapid iteration from concept to pilot.

Ideate | Evaluate | Prioritise | Solution Design | Build & Test | Deploy | Scale & Optimise

Cross-functional pods operating within established COE model and program structure

Role/Function	Responsibilities in AI Delivery		Why
Business Product Owner (e.g., Head of Claims, UW, or Policy Admin)	R	Prioritizes features based on business value	Anchors solution in real insurance processes (FNOL, premium calc, endorsements)
	A	Validates user stories and functional workflows	
	A	Acts as bridge between IT and operations	
Domain SMEs (Claims Analysts, Underwriters, Actuaries)	C	Provide business rules for rating, reserving, claims triage	Define how system should behave for each LOB (Line of Business)
	C	Verification of business rules and workflow processes	
	C	Participate in gap analysis and UAT	
CIO Organization (Platform Teams) COE Model	R	Pioneer technology introduction and pilot with business	RSA's established model for tech delivery at scale
	A	Delivers common capabilities (Copilot, Instabase, Power Platform)	
	I	Provides technical architecture and integration patterns	
Data & Analytics Team (Hub-and-Spoke Model)	R	EDP integration, data quality frameworks, analytics infrastructure	Central platform + federated BI/MI ensures data foundation
	A	Data scientists embedded in business teams (Claims, Pricing)	
	C	ML model development and validation	
Enterprise Architecture (Graham Wilde & Team)	A	Solution architecture, integration patterns, technical feasibility	Ensures AI integrates with modernization roadmap (decommission legacy by 2028)
	A	Alignment to Architecture Blueprint and strategic platforms	
	C	Avoid tactical "Unicorn" solutions and technical debt	
Ethical SteerCo + Working Group (Andrew Tongue & AI Council)	A	Approve guiding principles for ethical AI use	FCA/GDPR/EU AI Act compliance embedded from inception
	A	Review/challenge AI initiatives for customer outcomes	
	R	Mandatory DPIA sign-off before production	

R Responsible A Accountable C Consulted I Informed

Operationalizing the TOM

Stage 4: Deliver — Extending Across Governance and Adoption (Part 2)

HEXAWARE

Delivery extends beyond build and test—embedding governance, change management, and user engagement to ensure every AI solution is ethical, integrated, and adopted at scale.

Ideate | Evaluate | Prioritise | Solution Design | Build & Test | Deploy | Scale & Optimise

Cross-functional pods operating within established COE model and program structure

Role/Function	Responsibilities in AI Delivery	Why
IT & Data Governance Leads	R Integration with core systems (PAS, CRM, Guidewire) A Infrastructure readiness (cloud/on-prem) A Data privacy, access controls, retention policies	Insurance deals with legacy tech, complex integrations, regulated data (PII)
Change Management & Training	R Prepare business users for new tools and workflows R Create role-based training (agents, brokers, adjusters) I Drive adoption through workshops and feedback loops	Front-line adoption by distributed users critical to success
End Users (Agents, Brokers, Claims Adjusters)	C Participate in design validation and usability testing C Report pain points and suggest improvements I Pilot rollouts and provide feedback	Primary system users - early feedback ensures adoption and reduces post-launch changes

Program Delivery Structure

- **Strategic Programs Own Use Cases** → CL Tech & Workflow, CL Pass, Guidewire, Integrated Platform Pricing, One Commercial, Symphony integration
- **Cross-Functional Pods** → 2-week Agile sprints with Product Owner + Scrum Master + Dev team + QA + Business SMEs
- **Governance Cadence** → Monthly steering committee + Quarterly gate reviews + Six-monthly Ethical SteerCo reporting

R Responsible | A Accountable | C Consulted | I Informed

HEXAWARE

Seven-Phase Journey for Use Case: Email/Document Triage for Mailbox (1/2)

End-to-end demonstration of how RSA's priority use case flows through all seven phases of the TOM

Use Case: Email/Document Triage for Mailbox

Quick Win

Business Problem

Brokers email submissions to individual underwriters (e.g., "Fred") instead of using group mailbox. Manual routing creates delays in triage. Inconsistent submission ingestion into Instabase AI Hub for document extraction. Lost efficiency in underwriting workflow.

Proposed AI Solution

AI classification model identifies broker submissions in individual Outlook inboxes, classifies submission intent/type (NB, MTA, Inquiry) , recommends routing to Dynamics CRM group mailbox. Safety Control: Human-in-the-loop approval (Segment 1 launch),-> High- confidence auto-routing (Segment 2 at >95% accuracy).

Downstream Integration: Auto-routed submissions trigger Instabase AI Hub ingestion for document extraction

Value Framework Scores

Business Impact: 3.6/5
(Revenue 3, Cost Savings 4, Risk 3, Broker Exp 4, Strategic Fit 4)

Implementation Effort: 2.2/5
(Data 2, Tech 2, Change 2, Model Risk 2, Adoption 3)

QUADRANT: Quick Win (High Impact 3.6, Low Effort 2.2)
Impact Category: Expense Ratio (UW Productivity)

T-Shirt Size

Size: S (Small) 

Timeline: 12-14 weeks | Team: 2-4 people | Investment: ~£265K

Projected Benefit: £500K-£1.2M annually | ROI: 2-4x

Note: "Segments" refer to functional implementation stages for the Use Cases; "Phases" refer to TOM delivery stages (all use cases)

HEXAWARE

Seven-Phase Journey for Use Case: Email/Document Triage for Mailbox (2/2)

End-to-end demonstration of how RSA's priority use case flows through all seven phases of the TOM

Ideate	Evaluate	Prioritise	Solution Design	Build & Test	Deploy	Scale & Optimise
Inside-Out: Graham Wilde: "Brokers email Fred directly, not group mailbox" Pain: Manual routing inefficiency Strategic: One Commercial UW efficiency **Outside-In:** Intact: Email classification deployed Hexaware: Email triage horizontal capability	**10-Lever Scoring:** High: Cost savings (manual effort), Broker experience (speed) Low: Tech complexity (integration exists), Data (metadata available) Moderate: Adoption (UW trust) **Governance:** DPIA initiated (email processing) Approval loop mitigates risk	**Quadrant:** Quick Win (3.6 impact / 2.2 effort) Q1 2026 deployment **Resources:** UW Product Owner, Power Platform dev, BA, QA, 5 SME underwriters Budget: ~£265K Timeline: 12-14 weeks **Dependencies:** Dynamics mailbox, Outlook API	**Architecture:** Power Automate flow (Outlook → AI → Dynamics) Approval UI design for underwriters Instabase trigger design **User Stories:** As UW, I want submissions auto-routed so I don't miss emails **EA Review:** Graham Wilde team validates Outlook-Dynamics integration pattern	**3 Sprints(Total 6 Weeks):** Build Power Automate flow Train NLP model on historical emails Develop approval UI **UAT (2 weeks):** UAT with 10 underwriters Feedback loop adjustments **DPIA:** Ethical Working Group sign-off (email content processing approved)	**Pilot Deployment:** For 2 weeks, Pilot with 10 underwriters Hypercare support (daily war room) Performance monitoring dashboard **Training:** User guides created Quick reference cards **Cutover:** Final DPIA confirmation Change Control Board approval	**Months 4-6:** Rollout to all Commercial Lines UW Tune confidence threshold Document integration pattern **Months 6-12:** Extend to Specialty Lines Replicate to Claims email triage Remove approval loop (full automation) **Lessons:** NLP model → horizontal capability Actual ROI vs. projection trackingX

Cross Functional POD structure for Use Case #3: Email/Document Triage for Underwriter Mailbox

Business Product Owner	CIO Scrum Team	Domain SMEs	Enterprise Architecture	Ethical Working Group	Change Management
Head of Commercial Lines Underwriting **Role:** Define priority, validate workflow, UAT sign-off	Power Platform dev, BA, QA engineer **Role:** Build Power Automate + NLP integration	5 Commercial Lines Underwriters **Role:** Test approval loop, validate classification	Graham Wilde EA team **Role:** Review Outlook-Dynamics integration (Phase 4)	Andrew Tong Governance **Role:** DPIA initiation (Phase 2), sign-off (Phase 5)	Training Lead **Role:** User guides and training (Phase 6)

HEXAWARE

Next Steps to Operationalize the TOM

Immediate actions, and strategic considerations for operationalizing the complete AI Target Operating Model

Immediate Actions Required (Next 30 Days)

Priority	Action	Owner	Success Criteria
P0	**DPIA Backlog Clearance Plan** Prioritized plan to complete DPIAs for 46 outstanding initiatives	Ethical Working Group Chair + Andrew Tong	Prioritized list with timeline; dedicated resources; streamlined template created
P0	**Operationalize Value Framework** Lock down 10-lever parameters; deploy prioritization tool	Tricia Deane (CIO) + Program Steering Committees	Agreed lever definitions; tool deployed; training completed
P1	**Quick Win Portfolio Activation** Select 3-5 Quick Wins (including Use Case #7) for Q1 2025	Business Product Owners + CIO Platform Teams	Pods formed; sprint planning completed; DPIA initiated; delivery kickoff
P1	**Cross-Department Alignment** Break down silos through unified 7-phase TOM adoption	Toby Harris (CIO) + Data Team + Graham Wilde (EA)	Joint working sessions; unified intake process; common framework adopted
P2	**MLOps Maturity Roadmap** Address gaps in model drift monitoring, retraining governance	Data & Analytics Team + CIO Platform Teams	MLOps maturity assessment completed; roadmap defined; tooling selected

Strategic Considerations: Intact Benchmarking & 2026 Targets

7-Phase Roadmap to Intact-Level Capabilities

Current State (2025): Use Case #3 (email routing) = foundational Phase 1-6 execution building organizational muscle

2026 Target: Progress to 75% submission pre-fill using LLM inference (handwritten/free-form text) with drag-and-drop broker uploads – matching Intact Canada 2024 achievement

7-Phase Strategic Roadmap:

Phase 1-3	Phase 4	Phase 5	Phase 6	Phase 7
Use Case identified, evaluated, prioritized	Enhanced Instabase AI Hub integration architecture designed	Full submission pre-fill with Power UP workbench built & tested	Data enrichment with risk hazards / firmographics deployed	Enterprise-wide scale achieving 75% pre-fill, continuous optimization

Expected Impact

- $150M+ recurring AI benefits by 2030 (Intact benchmark)
- 4.7pts CR outperformance through AI-driven pricing/risk selection

Use Case Portfolio & Roadmap – Prioritized Quick Wins and Path to Progression

This section presents RSA's **four prioritized quick-win use cases** — each validated with stakeholders for value and feasibility — and outlines a **directional roadmap** that connects these early initiatives to RSA's longer-term AI progression and capability building.

Note: A complete backlog of 50+ use cases, including CIO submissions, will accompany this report as a separate Excel annex.

Use Case Portfolio & Roadmap

AI Technology Vision for 2030

RSA's near-term use cases are the first steps toward this evolution — shifting from people-driven to agentic AI–enabled services.

People driven Services	Technology Driven Services	Platform Driven Services	AI – Led Agentic Services	Agentic Ecosystem driven services
Consultants providing managed services with ability to scale up or down. Ramp up and Ramp down happens based on projects	*Primarily driven by people but supported by proprietary solution accelerators, tools, and software.*	*Leverage built-in delivery platforms to enhance service delivery and efficiency. People support for customizations*	*Augmenting human capabilities with smart AI agents to optimize business processes and decision-making.*	*Focus on delivering services primarily through technology, minimizing human intervention, and maximizing efficiency.*
Key Features: • **Flexibility**: Easily adjust team size based on project needs. • **Expertise:** Access specialized skills not available in-house. • **Control:** Maintain direct oversight of projects and processes.	**Key Features:** • **Human-Centric:** Primarily driven by skilled professionals. • **Tool-Supported:** Utilizes a variety of technology tools and accelerators. • **Efficient:** Enhances service delivery through tech integration	**Key Features:** • **Integrated Platforms:** Uses cohesive platforms for service delivery. • **Scalability:** Easily scalable and consistent across various operations. • **Efficiency:** Enhances productivity and efficiency through platform support.	**Key Features:** • **AI-Augmented:** Combines human expertise with AI agents. • **Cost-Effective:** Achieves lower TCO through optimization. • **Enhanced Capabilities:** Expands service potential with AI-driven insights.	**Key Features:** • **Technology-driven:** Primarily led by advanced software solutions. • **Minimal Human Intervention:** Reduces reliance on human resources. • **Efficient and Scalable:** Provides efficient, scalable, and consistent service delivery.

2023 – 2026 **2027 – 2030**

Human Dependency Technology Dependency

AI Use Case Classification Model

HEXAWARE

RSA's near-term use cases are the first steps toward this evolution — shifting from people-driven to agentic AI–enabled services.

	Level 0	Level 1	Level 2	Level 3	Level 4
Foundational Characteristics	**Rule Based Automation** Predefined, rules-driven workflows; no learning or adaptation	**Assistance (Retrieval)** Retrieves relevant information from structured/unstructured sources	**Augmentation (Summarization)** Summarizes, contextualizes, and links insights across sources utilising multiple tools	**Collaboration (Orchestration)** Coordinates multiple agents; executes workflows across domains.	**Autonomy (Self-Management)** Self-directed, goal-oriented agents that plan, reason, and act with minimal supervision
	Algorithmic automation	**Foundation models and basic agents**	**Advanced agents with memory and reasoning**	**Agentic systems**	**Autonomous systems, finetuned models**
			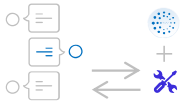	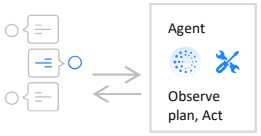	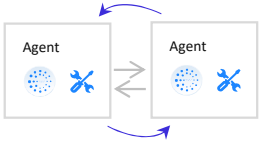
Technical Components	BPM engines, RPA bots, Legacy rate tables, Rule engines	NLP chatbots, OCR, Vector DBs, RAG systems, Stateless execution	API integrations, Knowledge graphs, Multi-source data integration platform, Agent Memory	Orchestration platforms, Microservices, API-first, event-driven, Integrated with EDP	Meta-learning systems, Autonomous orchestration, Self-optimizing, Domain SLM

Rigid, Reactive, Static ← *AI Assistants* — **Agentic AI is a continuum** — *Agentic Workflows* → **Proactive, Independent, Evolving**

The more adaptive the system, the more complex the oversight needed (human-in-the-loop validation, continuous monitoring, robust AI governance)

AS2

Beyond Quick Wins – RSA's Broader Use Case Landscape

HEXAWARE

RSA's AI journey extends beyond initial quick wins, drawing from a diverse pipeline of more than 50 ideas that span functions and maturity levels to sustain innovation momentum.

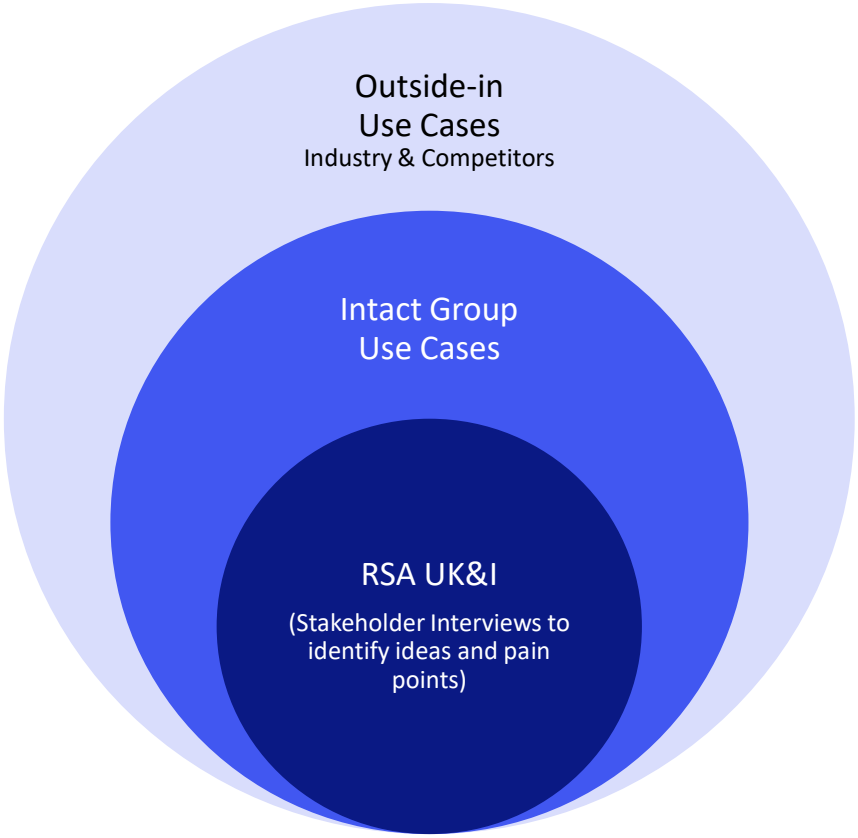

DP1

- **Current backlog**: ~50+ ideas logged.
- **Candidate pool**: ~20 shortlisted through Hexaware/RSA filtering.
- **Validated quick wins**: 2–3 ready to execute now.

- This landscape becomes the pipeline for RSA's AI learning agenda.

Slide 38

DP1 Update graphical representation with this context.
Donna Pahel, 2025-09-26T12:08:31.968

ST1 0 The numbers are present in the upcoming slide
Shreyash Tiwari, 2025-09-26T12:12:00.196

DP1 1 [@Shreyash Tiwari] please update to make relevant and correct. Thank you.
Donna Pahel, 2025-10-07T21:40:59.952

AS2 Category inconsistent in this block/part of the presentation
Anna Shymchenko, 2025-10-02T11:43:21.925

Use Case Portfolio & Roadmap

AI Transformations Across the Insurance Value Chain

HEXAWARE

RSA's prioritized use cases map to foundational enablers across the insurance value chain—advancing AI integration in Claims, Underwriting, and Distribution while establishing reusable, cross-enterprise capabilities.

Risk & Data Acquisition	Underwriting	Distribution & Delivery	Marketing & Sales	Claims & Customer Support
Automated KYC and identity verification	Underwriting Decision Agent	Personalized onboarding journeys for customers	Hyper-personalized campaign content generation	Claims triage automation (first notice of loss to categorization)
AI-driven fraud detection & risk scoring	Automated decision support from large document ingestion	AI-powered agent/advisor support during policy sale	Lead prioritization & conversion assistant	Intelligent claims settlement assistant
Smart data enrichment from external sources	Risk-based premium recommendation engine	Workflow automation for faster policy issuance	AI-based cross-sell and upsell recommendations	GenAI chat assistant for customer queries

Domain Use Cases

These Specialized Use Cases are designed to address business processes in each domain, powered by reusable system agents.

Foundational Use Cases

These are reusable, modular, interoperable code components that can be redeployed throughout the enterprises

Foundational Use Cases / Capabilities

Data & Recognition
- OCR Agent
- Data Ingestion Agent
- Data Validation Agent

Knowledge & Content
- Pre-processing Agent
- Summarizing Agent
- Logging Agent

Automation & Monitoring
- Scheduling Agent
- Notification Formatter Agent
- Monitoring Agent

Learning & Governance
- Risk Learning Agent
- Supervisory Agent
- Test Case Execution Agent

HEXAWARE

Global View of Intact Group Use Cases

The Intact Group portfolio provides proven use cases across the insurance value chain—offering validated models, shared learnings, and governance practices that inform RSA's AI roadmap.

Value chain	Concrete Data & AI use cases	Status/Timeline	Where & focus
Distribution & Broker	AI-enhanced omnichannel (client agent matching), speech analytics (call reason detection); digital FNOL & tracking; high adoption of mobile/app	Deployed (portfolio since 2018); ongoing scale	Group; Canada leads in digital, UK& adopting under modernization/One Commercial.
Submission Intake	Automated quote parsing & pre-fill (LLM inference, handwriting/free- form; drag-and-drop broker uploads)	Live in Canada (2024) →Expand US & UK; CL/SL pre-fill scaling by 2026	Group program; critical for GSL (speed-to-quote) and UK&I (sales productivity).
Triage & Prioritization	AI triage for claims surge (geospatial mapping, sophisticated triage); emerging trend & high-exposure identification	Deployed; ongoing enhancement	Group; GSL/UK&I benefit in CAT and complex claims.
Data Enrichment	Enrich quotes with risk hazards, firmographics, nearby risky ops, weather, crime	Live in Canada (2024) -> Extend US & UK; more data points by 2026	Group; GSL specialty risk insights; UK&I segmentation gains.
Risk Assessment	ML risk/profit models; 10x LR separation example (Canada Comm. Property)	Deployed; continuous model refresh	Group; principles exported to UK&I/GSL with governance.
Quantitative Underwriting	Quantitative UW frameworks; governance of deviation from modeled price at per-risk level	Deployed; embedded governance	Group; crucial for London/GSL & UK&I to hold CR targets.
Pricing	Optimized price strategy via ML & elasticity; strict per-risk governance	Deployed; widening coverage of DPW using ML (PL, CL/SL; Canada, UK, GSL)	Group; scale to UK&I/GSL part of 2026 agenda.
Policy Servicing	AI in operations: document summarization, speech analytics/coaching; digital service features scaling	Deployed; expanding	Group; UK&I modernization calls out data/AI uplift to service & productivity.
Claims FNOL-> Settlement	Fraud detection models, total loss determination & market value suggestion, one-touch settlement; pipeline for automated assessment, vendor recommendation, automate repairables	Many deployed; pipeline items listed as "In the pipeline"	Group; GSL gains on complex lines; UK&I unlocking CR/ER via ops excellence.
Supply Chain & Repair	Owned repair network (On Side) + data-driven orchestration: 46% of property repairs; 30+ days faster; 33% faster cycle; +12 NPS	Deployed; scaled	Group advantage, exportable playbook.
Telematics & Climate	4th-gen UBI; 67-pt LR gap best vs. worst third; climate risk modeling for +3-5°C future	Deployed; maturing	Group; learnings flow to UK&I personal/commercial auto contexts and GSL where relevant.
Investments	Quantitative model for IIM (investment management	Deployed	Group portfolio alpha

HEXAWARE

Value Framework

A structured framework was applied to evaluate use cases by business value and implementation complexity—identifying opportunities that deliver measurable impact with minimal friction to deploy.

Output

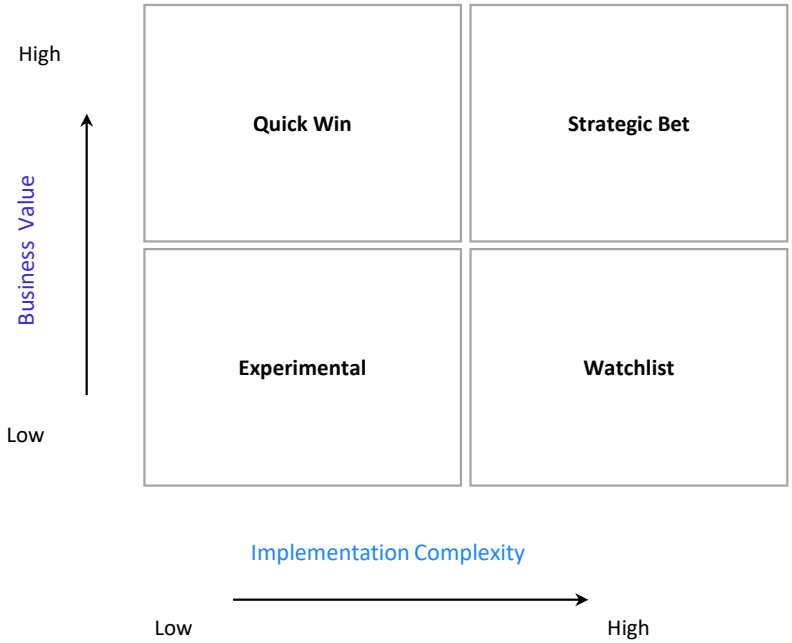

Executive Dashboard
Interactive matrix plots with quadrant-based ROI analysis

T-shirt Sizing System
UK benchmark-compliant cost & timeline estimation

Input

Matrix to focus on Implementation Complexity of Use Case vs Business Value.

Some of the key parameters considered are:

Business Impact Levers

- Revenue Impact
- Cost Savings
- Risk Reduction
- Broker Partner Experience
- Strategic Fit

Implementation Effort Levers

- Data Readiness
- Technical Complexity
- Change Impact
- Model Risk
- Adoption Readiness

- The framework uses a 1–5 scale for each lever. Higher values correspond to greater business impact or increased complexity.
- Each Lever is currently weighted at 20%

LEGO Architecture
Rational, reusable component with 60% maintenance reduction

User Feedback Integration
Continuous improvement through contextual feedback collection

HEXAWARE

Use Case Dashboard

Through structured evaluation and stakeholder validation, 51 logged ideas were assessed and segmented into high-, medium-, and quick-win priorities, focusing initial activation on four near-term opportunities.

51	24	25	4
Total Use Cases	High Priority Use Cases	Medium Priority Use Cases	Quick Wins Identified

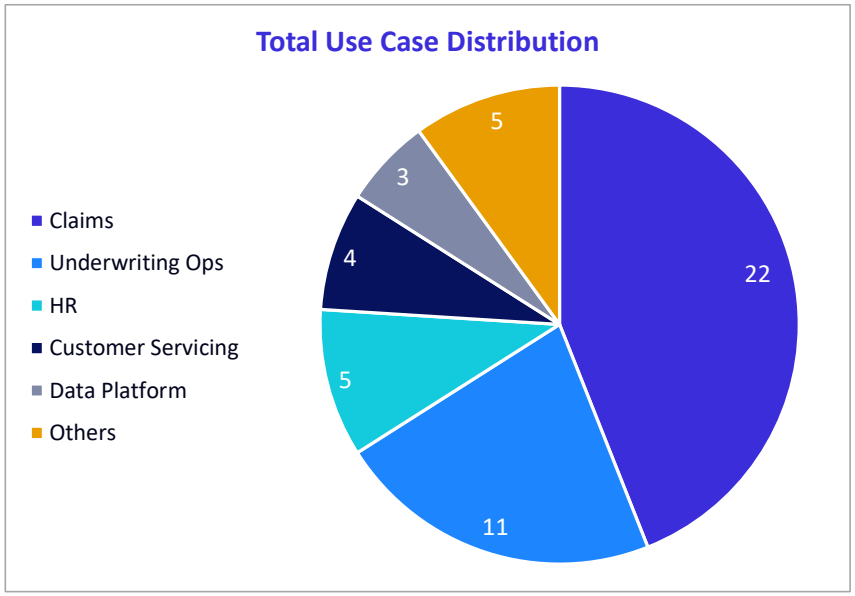

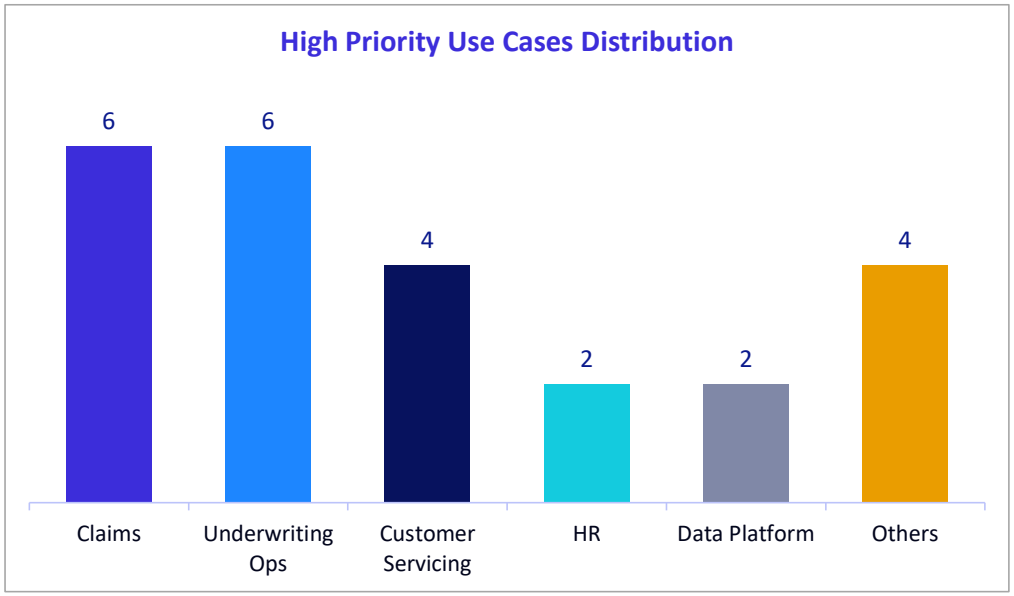

DP1

Use Case Portfolio & Roadmap

HEXAWARE

Quick Wins (High Impact / Low Effort Use Cases)

The four selected quick-win use cases represent immediate, high-confidence opportunities to demonstrate GenAI value, build internal confidence, and create reusable foundations for scale.

Email/document triage for mailbox

Reduces underwriting admin while testing how GenAI embeds into workflows

Underwriting Ops

Instabase ingestion validation & external data enrichment

Supports broker profiling, cross-sell, prospect insights. Reusable across multiple domains

Data & Risk Validation

Call summarization for claims interactions & sentiment analysis

Call summarisation and sentiment analytics for claims interactions

Claims

AI driven broker analytics for cross sell

Brokers lack proactive cross-sell/prospecting insights → AI-driven Broker Analytics

Sales & Distribution

Rationale for Selection

- **Quick technical win:** Low-code Power Platform leverages existing infrastructure
- **Validated by Graham Wilde**
- **Solves real pain point:** Brokers email individual underwriters instead of group mailbox
- **Trust-building approach:** Human-in-loop (Segment 1) → graduated autonomy (Segment 2 at >95% accuracy)
- **Scalable learning:** Reusable pattern across LOBs; extends to claims email triage
- **Impact:** ~0.15-0.2 CR pts | Investment: £265K | Timeline: 12-14 w

- **Extends proven capability:** Instabase deployed for Commercial Combined; validation adds immediate value
- **Two-segment approach:** Segment 1 (Q1 2026 validation) + Segment 2 (Q2-Q3 2026 enrichment) = manageable risk
- **Automates compliance:** 100% sanctions/watchlist checks + 80-90% data validation automation
- **Proposed by Andrew Hall** supported by underwriters
- **Foundation for scale:** Reusable validation layer across all underwriting & claims LOBs
- **Impact:**~0.2-0.3 CR pts | Investment:1.35- 1.75M | Timeline: 16-18 w (two segments)

- **Infra accelerator:** Call Miner licensed & integrated with Vodafone/Cisco telephony
- **Validated productivity gain:** James Gregory confirmed 6-7 FTE improvement (£300-400K annual savings)
- **Project momentum:** Mobilization already underway with assigned PM from Claims team
- **Dual-segment value**: Segment 1 (end-of-call summary) + Segment 2 (live agent assist prompts)
- **Reduces manual burden:** Eliminates inconsistent note-taking, improves audit trail
- **Impact:** ~0.15-0.2 CR pts | Investment: £1.0M | Timeline: 12-14 w

- **Strategic priority: Validated Quick win**
- **Executive engagement model:** CEO/Regional Director escalation framework (Ken + Tova alerts)
- **Leverages existing assets:** 19+ broker portals, Unicorn database, Qlik dashboards already in place
- **High revenue impact:** 10-12% retention improvement + 15-20% faster pipeline conversion
- **Low incremental cost:** Uses existing infrastructure (Unicorn & Qlik)
- **Impact:** ~0.1-0.15 CR pts | Investment: £1.2M | Timeline: 14-16 w

Slide 43

DP1 [@Mukundan Madhavan] and [@Anton Makarevich] - please check content for accuracy.
Donna Pahel, 2025-10-08T00:31:31.504

AM1

Use Case 1 : Call summarization for claims interactions & sentiment analysis (1/4)

HEXAWARE

Problem Statement	Claim handlers face operational inefficiencies and elevated error risk due to manual extraction and documentation of critical information from unstructured call transcripts.
Proposed by	Vicky Walker (UK); acknowledged by Andy Flower. Further validated by James Gregory
Current State	**Extracting context:** • Currently all claims handlers read through long, often unstructured transcripts to identify key details(incident description, policy coverage, claim status, claimant emotions, etc which is time consuming and inefficient. • Urgent details maybe buried deep in the transcripts. • Call Miner Platform with MS Cognitive Service is licensed and integrated to telephony(Vodafone/Cisco). The Platform includes desktop client capability for reali time transcription and live prompting, but **this functionality is not mobilized for production use.** • Current 2 hour lag exists due to Verint call extractor batch processing. **Note taking:** • Manual notes are subjective, shorthand heavy, and vary greatly by handler. • This creates inconsistency in notes. In multi touch points , subsequent handlers may spend extra time decoding the previous notes leading to duplication of work. Minimal audit trail. **Case progression:** • Delays in claim progression due to lag between extraction , reading and documentation as all these are currently executed manually.
Use Case / Automation Potential	• Automated Transcript Summerisation • Intelligent Information Extraction • Automated Note generation • Workflow integration and Alerts
Systems Impacted	• Guidewire ClaimCenter notes, Instabase, legacy doc repositories, MI/reporting platforms. • Existing call-transcription platform (Call Miner); Guidewire (claims notes/communications

Slide 44

AM1 my understanding from the last call with James Gregory is that the next 6 slides (2 cases) are exactly the same project that is already ongoing. information just came to us from 2 different sources (Vicky and James himself). But Vicky was also mentioned as a supervisor of the Jame's project

Anton Makarevich, 2025-09-29T11:06:11.468

AM1

HEXAWARE

Use Case 1 : Call summarization for claims interactions & sentiment analysis (2/4)

Assumptions & Dependency	CallMiner, Microsoft Cognitive Services, and Verint are already procured and integrated with telephony systems Vodafone/Cisco telephony should have “real-time” integration with CallMiner The project mobilization is ongoing with assigned PM from Claims team
Input Required	Live call audio streams from telephony system; historical call data for AI model training/calibration; compliance rules and prompting logic; agent desktop integration endpoints; CRM data for context-aware suggestions.
Target/ Future State	**SEGMENT 1: IMMEDIATE (Real-Time Transcription & Summary)** • Real-time call transcripts with AI-generated end-of-call summaries • Immediate AI summary and action capture at wrap-up • Automated quality and compliance checking • Real-time alerts and coaching signals: Supervisors receive immediate notifications for quality/compliance issues, enabling proactive intervention **SEGMENT 2: POST-STABILIZATION (Live Agent Assist)** • AI provides live compliance prompts during the call • Agent assist suggestions during active calls • Live transcription of conversation as it happens via desktop client **SYSTEM INTEGRATION:** • Guidewire ClaimCenter, CallMiner (licensed), Vodafone/Cisco telephony, Desktop client **MANUAL OVERSIGHT:** • Supervisor review of flagged issues; human validation for high-risk claims
Technology & Data Considerations	• Real-time transcription latency must be <2 seconds for effective agent assist • System must handle high concurrent call volumes without degradation • Data privacy and compliance requirements for call recordings • Integration between systems must be stable and performant
Output	• **Segment 1 (Immediate):** Real-time call transcripts with AI-generated end-of-call summaries and compliance alerts. • **Segment 2 (Post-stabilization):** Live agent assist prompts during calls via desktop client. **Expected Benefits:** 300-400k Annual savings ~6-7 FTE productivity improvement (validated by James Gregory)

Note: "Segments" refer to functional implementation stages for the Use Cases; "Phases" refer to TOM delivery stages (all use cases)

Slide 45

AM1 updated
Anton Makarevich, 2025-09-29T10:57:04.444

AM3

Use Case 1 : Call summarization for claims interactions & sentiment analysis (3/4)

HEXAWARE

ST2

Business Value KPIs

BASELINE METRICS:
- Manual note-taking: Subjective, shorthand-heavy, inconsistent across handlers
- 2-hour lag due to Verint batch processing
- Claims handlers read through long unstructured transcripts manually
- Multi-touch points: Subsequent handlers spend extra time decoding previous notes
- Minimal audit trail; duplication of work common

TARGET METRICS (12–18-month realization):
- £300-400K annual savings
- 6-7 FTE productivity improvement
- Real-time transcription (<2 second latency)
- 100% call coverage with automated quality checking
- Supervisor alerts for quality/compliance issues (proactive intervention)

BUSINESS OUTCOMES:
- ✓ Reduced handling time, faster updates, improved note consistency & auditability
- ✓ Benefits span customers, brokers, and staff—lower expenses, smoother journey
- ✓ Frees up team time to focus on complex claims by eliminating non-value tasks

Risk & Mitigation

Risk: Inaccurate or incomplete call summaries may lead to incorrect claim decisions or miscommunication.
Mitigation: Implement human-in-the-loop review for high-risk claims and validate AI outputs against key action items.

Implementation Timelines with Standard roles requirements

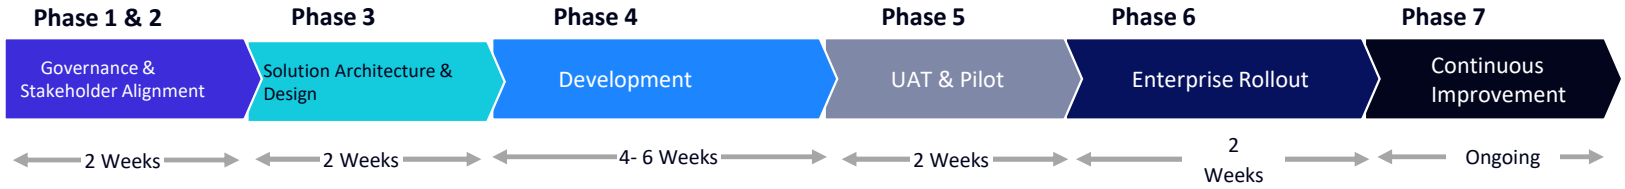

Total : 12-14 Weeks

Roles:

Data scientist/ML engineer (0.5), Backend developer (2), DevOps engineer (1), Business Analyst (0.5), Solution/Technical Architect (1), QA engineer (1)

High Level Cost Estimation

T-Shirt Size: MEDIUM (justified by existing CallMiner infrastructure)

- Build Cost: £650-750K
- Licensing Cost: £0 (CallMiner licensed, MS Cognitive Services included in EA)
- Annual Run Cost: £150-200K
- Infrastructure Cost: £100-150K

ST1

Total Q1 2026 Investment: ~£1.0M
- **Cost Breakdown by Phase**:
  - Phases 1-3 (Governance): £50K (DPIA, stakeholder workshops)
  - Phase 4 (Architecture): £75K (integration design, EA validation)
  - Phase 5 (Development - 3 sprints): £400-500K (real-time integration, NLP model, desktop client)
  - Phase 6 (UAT/Pilot): £100K (10 Claims handlers, pilot monitoring)
  - Phase 7 (Rollout): £275-325K (enterprise deployment, model retraining pipeline)
  - - Contingency (15%): £100K

Slide 46

ST1 [@Shruti Lawani] [@Mukundan Madhavan] Please add T-Shirt Sizing
Shreyash Tiwari, 2025-09-29T14:57:37.828

ST2 [@Anton Makarevich] Please provide rough estimate of roles & timelines
Shreyash Tiwari, 2025-09-30T08:45:26.521

AM2 0 the team is always similar, for the timeline, this onle is smaller than the others - I would go with 6 sprints/12 weeks, not 8/16 as other use cases. including preparation and stabilization sprints, so 4 to actually build the pilot
Anton Makarevich, 2025-09-30T11:49:08.138

AM3 that should be aligned with the diagram we have in TOM now
Anton Makarevich, 2025-10-06T08:55:20.029

AM3 0 but the challenge is that the TOM is a next section now, still i believe we should be consistent
Anton Makarevich, 2025-10-06T08:56:46.351

Use Case 1 Implementation Plan: Call summarization for claims interactions & sentiment analysis (4/4)

HEXAWARE

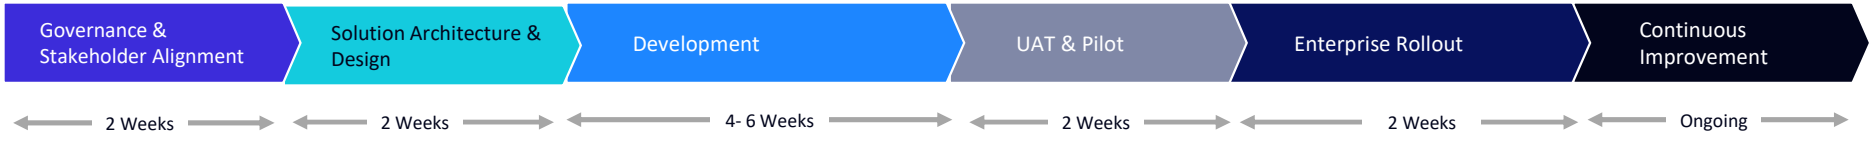

	Phase 1 & 2	Phase 3	Phase 4	Phase 5	Phase 6	Phase 7
Activities	• DPIA completion (already in progress • Stakeholder sign-off • Benefits validation (£300-400K, 6-7 FTE)	• CallMiner platform mobilization plan • Real-time transcription architecture validation • Compliance prompt logic definition • Desktop client integration design	• **Sprint 1 (2 weeks):** Real-time CallMiner integration + Guidewire notes API • **Sprint 2 (2 weeks):** AI summarization model + compliance alerting • **Sprint 3 (2 weeks):** Desktop client prompts + quality metric	• UAT with 10 Claims handlers • Pilot deployment in Claims + E-Trade • Performance monitoring (<2 second latency requirement) • Feedback collection	• Phased rollout to all Claims teams • Supervisor alert dashboard • Performance optimization	• Performance monitoring • Accuracy improvement iterations • Phase 2 planning (live agent assist expansion)
Key Deliverables	• Approved DPIA • Executive sponsor commitment • Success metrics defined	• Validated architecture • Integration endpoints documented • Compliance rules catalogued	• Real-time transcription enabled • End-of-call summary generation • Compliance alert system • Agent assist UI (desktop client)	• 10 Claims handlers trained • Pilot metrics dashboard • Issue log and resolution tracking	• Enterprise-wide deployment • Supervisor coaching signals operational	• Monthly model performance reports • Phase 2 roadmap
Critical Success Factors	• DPIA already in progress (accelerator) • Strong stakeholder validation exists	• Architecture documented, requires validation only (accelerator) • CallMiner platform already integrated	• Vodafone/Verint contract for real-time audio streaming **(Critical)** • Microsoft Cognitive Services API integration	• Real concurrent call volume testing • Handler feedback on agent assist prompts	• Change management for all handlers • Telephony system stability under load	• Feedback loop with Claims leadership • Continuous compliance rule updates

Use Case 2 : AI-Driven Broker Analytics for Cross Sell (1/4)

HEXAWARE

Problem Statement	Brokers operate without real-time cross-sell or prospecting intelligence, limiting strategic growth opportunities.
Proposed by	James (Jay) Gregory (walked through Unicorn, portal data, and target state for automated insights).
Current State	**Pipeline Generation** • Pipeline tracking is currently manual, relying on inputs from sales teams, underwriters, and Excel dumps from brokers or online sources. • RSA has deployed Qlik dashboards with broker facing story packs that track inquiry-to-quote conversion, product splits and performance metrics. Unicorn SQL database contains structured broker data including performance history, but lacks automated propensity scoring, opportunity prioritization, and executive alerting mechanisms. Currently 19+ strategic broker portals (including Marsh, Aon) are manually scraped for pipeline data. • Data from 20 Strategic Brokers is collected from their websites and analyzed over time but lacks real-time insights. **Growth** • The current pipeline holds untapped potential due to rich data sources from brokers and prospecting portals. • There's a clear opportunity to enhance quote volume by translating this data into a more focused and actionable pipeline. **Distribution** • Broker engagement is supported by manual scraping of portals and ad-hoc preparation for meetings. • There is no system in place for propensity scoring or automated opportunity surfacing. • Distribution efforts are reactive rather than predictive, limiting efficiency and scalability.
Use Case / Automation Potential	• An AI agent could proactively surface relevant profiles for quoting, reducing lag and improving targeting precision. • Leveraging AI could significantly accelerate identification of quotable prospects and streamline growth efforts.
Systems Impacted	Unicorn data model (SQL outputs, controls), broker data feeds (Marsh, Aon portals), Qlik/Power BI, CRM (MSD) for execution.

AM1

Use Case 2: AI-Driven Broker Analytics for Cross Sell (2/4)

HEXAWARE

Assumptions & Dependency	Broker portals access/authentication is well documented, data can be normalised Unicorn model is documented, schemas are available Foundational cloud (Azure) infrastructure (including event mesh to support event-driven architecture) already exists
Input Required	Brokers data, including Account/Policy information and opportunities information Existing internal data (CRM/claims system, Unicorn database), historical win/loss data

Target/ Future State	**AUTOMATED DATA INGESTION (Event-Driven or Scheduled)**
	• AI bot automatically ingests data from 19+ broker portals (Marsh, Aon, etc.) into Unicorn SQL database
	• Eliminates manual portal scraping and Excel data dump processing
	INTELLIGENT ANALYSIS ENGINE
	• **AI scoring engine** evaluates broker opportunities with:
	  • **Propensity to win scoring**
	  • **Cross-sell likelihood assessment**
	  • **Value/revenue potential ranking**
	• - Results stored in CRM for immediate access
	INTELLIGENT REPORT GENERATION & EXECUTIVE TRIAGE
	• **Automated** broker-facing reports (monthly cadence)
	• **Risk-rated escalation framework:**
	  • **Large strategic brokers (declining metrics) → Auto alert to Ken (CEO) for immediate intervention**
	  • **Medium-tier brokers (declining metrics) → Alert to Regional Directors (e.g., Tova)**
	  • **Small brokers → Automated outreach via sales team**
	• **AI-generated opportunity insights embedded in Qlik dashboards**
	• **Integration with Unicorn Database for historical performance context**

Technology & Data Considerations	• Message broker: Azure service Bus, Azure Event Hubs, Kafka (managed or Confluent) • Compute: AKS, Azure Functions • AI model and platform: AI Foundry + Azure OpenAI • ML model, algorithms and platform: Azure ML, Databricks • Analytics dashboards: Power BI
Output	Broker specific reports, scored opportunities pipeline, Power BI analytical dashboards, automated alerts and notifications

AM1 updated
Anton Makarevich, 2025-09-29T09:11:57.129

Use Case 2: AI-Driven Broker Analytics for Cross Sell (3/4)

HEXAWARE

Business Value KPIs

BASELINE METRICS:
- 19+ broker portals serving ~£2.7B DPW
- Manual monthly reporting cycle (portal scraping Excel trackers)
- No propensity scoring or automated opportunity surfacing
- Reactive distribution approach limiting growth efficiency

TARGET METRICS (12–18-month realization):
- 10-12% retention improvement via proactive broker engagement
- 15-20% faster pipeline-to-quote conversion
- Automated monthly broker reports (first-of-month cadence)
- CEO/Regional Director escalation framework operational

BUSINESS OUTCOMES:
- ✓ Higher win rates via prioritized outreach
- ✓ Shorter prep time; consistent data-driven broker engagement
- ✓ Stronger broker relationships and revenue growth
- ✓ Executive visibility across RSA/NIG portfolios

Risk & Mitigation

- **Risk:** Broker portal data may be inconsistent, leading to unreliable AI scoring & hesitation from underwriters to trust the outputs
- **Mitigation**: Implement strict data quality checks during ingestion & equip sales teams with meeting-ready reports to build confidence

Implementation Timelines with Standard roles requirements

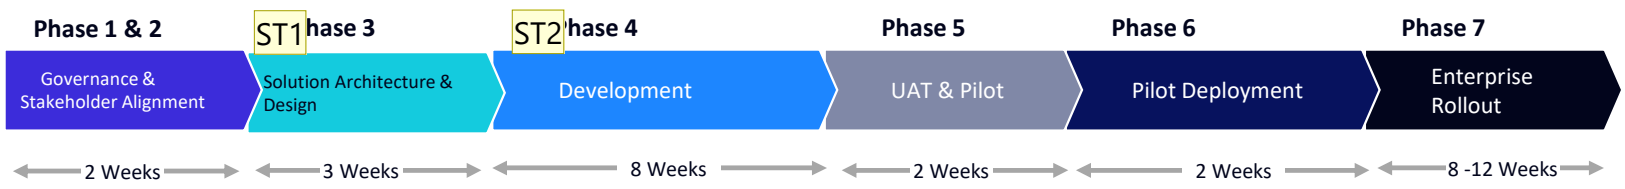

Total : 14 - 16 Weeks

Roles:
Data scientist/ML engineer (1), Backend developer (2), DevOps engineer (1), Power BI integration engineer (0.5) , Business Analyst (1), Solution/Technical Architect (1), QA engineer (1)

High Level Cost Estimation

T-Shirt Size: LARGE
- (19+ integrations, executive dashboard, 4-sprint development, multi-tier alerting)

Cost Breakdown by Category:
- Build Cost: £800-950K
- Licensing Cost: £50-100K
- Annual Run Cost: £150-200K
- Infrastructure Cost: £100-150K

Total Q1 2026 Investment: ~£1.2M
Cost Breakdown by Phase:
- Phases 1-3 (Governance): £50K (stakeholder workshops, prioritization framework)
- Phase 4 (Architecture): £100K (integration design, Unicorn schema analysis)
- Phase 5 (Development): £600-700K (4 sprints across multiple systems)
- Phase 6 (UAT/Pilot): £150K (5 strategic brokers, executive dashboard testing)
- Phase 7 (Rollout): £300-400K (19+ portal rollout, automated reporting, retraining)
- Contingency (15%): £150K

Slide 50

ST1 [@Shruti Lawani] & [@Mukundan Madhavan] Please add T-shirt Sizing to these
Shreyash Tiwari, 2025-09-29T15:00:20.604

ST2 [@Anton Makarevich] Please provide rough estimate of roles & timelines
Shreyash Tiwari, 2025-09-30T08:45:58.844

AM2 0 the same team + maybe extra Power BI specialist. standard 8 sprints/16 weeks for pilot
Anton Makarevich, 2025-09-30T11:51:30.929

Use Case 2 Implementation Plan: AI-Driven Broker Analytics for Cross Sell Timelines (4/4)

HEXAWARE

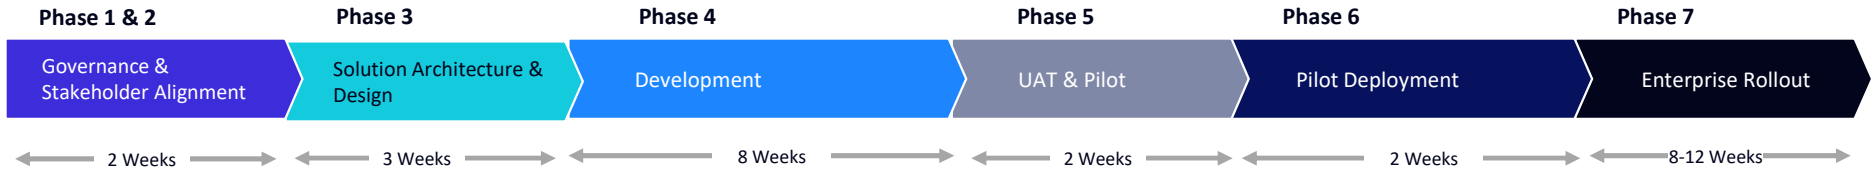

	Phase 1 & 2	Phase 3	Phase 4	Phase 5	Phase 6	Phase 7
Activities	• Use case validation with Kavit Rughani (Strategy) & James Gregory • Ken (CEO) + Tova (Regional Director) escalation framework approval	• Unicorn database integration design • Power BI dashboard enhancements • Propensity scoring model architecture • CRM (Microsoft Dynamics) integration pattern	• **Sprint 1 (2 weeks):** Unicorn DB integration + data pipeline • **Sprint 2 (2 weeks**): NLP model for opportunity scoring • **Sprint 3 (2 weeks):** Propensity model + CRM integration • **Sprint 4 (2 weeks):** Power BI dashboard + executive alerting	• UAT with 5 strategic brokers (Marsh, Aon included) • Executive dashboard testing (Ken, Tova) • Alert accuracy validation • Monthly report review	• Pilot with Marsh/Aon + 3 additional strategic brokers • Executive dashboard live for Ken • Regional director alerts enabled for Tova • Performance tracking	• Rollout to all 19+ broker portals • Automated monthly reporting enabled • Model retraining pipeline • Qlik dashboard story pack integration
Key Deliverables	• Executive sponsor commitment • Escalation authority matrix • Success KPIs (10-12% retention, 15-20% faster pipeline conversion)	• Unicorn database schema documentation • API access design • Dashboard wireframes • Scoring model framework	• Automated broker portal data ingestion (19+ portals) • Opportunity propensity scoring • Risk-rated escalation alerts • Monthly broker report generation	• 5 broker pilot validations • Executive feedback incorporated • Alert threshold tuning	• Pilot metrics dashboard • 5 strategic brokers monitored • Executive alerts operational	• All 19+ portals integrated • Automated monthly cadence • Self-service Qlik dashboards
Critical Success Factors	• Validated by stakeholders as "top priority quick win" (accelerator) • Clear executive engagement model	• Unicorn database well-documented • Qlik dashboard existing assets leveraged	• Broker portal authentication credentials (19+ portals) • Historical win/loss data quality • API rate limits on broker portals	• Strategic broker engagement • Executive availability for dashboard review	• Broker relationship impact monitoring • Ken/Tova feedback integration	• Broker portal authentication at scale • Data quality across all portals • Change management with sales teams

ST1

Use Case 3 : Email/Document Triage for Mailbox (1/4)

HEXAWARE

Problem Statement	Manual email triage of broker submissions causes delays, errors, and limits scalability.
Proposed by	Proposed enhancement by Hexaware; confirmed by Andrew Hall, Technical Validation: Graham Wilde (EA)_ “technically looks good….AI Model….Outlook integration….feedback loop with user approval”
Current State	**Process Inefficiencies** • Broker submissions (new business, MTAs, inquiries) arrive via email with no structured intake. • Brokers email individual underwriters (e.g: “Fred”) directly instead of using group mailbox • Underwriters manually review and forward relevant content to Instabase mailboxes. • Current deployment: Instabase active for Commercial Combined LOB only; other LOB’s lack automation • Lack of automation leads to delays, inconsistent triage, and missed opporti=unities **Human Dependency & Scalability Risks** • Workflow heavily relies on individual underwriters’ judgment and availability. • No standardized routing or prioritization logic across teams. • High risk of bottlenecks, errors, and knowledge silos as volume scales. **Operational & Data Integrity Challenges** • Manual handling increases chances of misrouting, duplication, or loss of critical documents. • Limited auditability and tracking of submission flow. • Inconsistent data capture affects downstream analytics and decision-making
Use Case / Automation Potential	**AI-powered email/document triage with graduated autonomy:** • **Segment 1 (Launch):** AI classification model identifies broker submissions in individual Outlook inboxes, classifies submission type (new business, MTA, inquiry), and recommends routing action with user approval loop before forwarding to Dynamics CRM group mailbox. • **Segment 2 (Post-validation):** Remove approval loop for high-confidence classifications (>95% confidence threshold) and enable full automation. • **Safety Control:** Human-in-the-loop validation initially to build underwriter trust and validate model accuracy before full automation. • **Downstream Integration:** Auto-routed submissions trigger Instabase AI Hub ingestion workflow for document extraction and data population.
Systems Impacted	**Primary Systems:** Outlook (individual underwriter inboxes), Dynamics CRM (group mailbox destination), Instabase AI Hub (downstream document processing) **Integration Layer:** Microsoft Power Platform (Power Automate) for workflow orchestration **API Dependencies:** Microsoft Graph API (Outlook integration), Dynamics 365 API (CRM routing), Instabase API (submission ingestion trigger)

Slide 52

ST1 [@Shruti Lawani] Please help us to fill this
Shreyash Tiwari, 2025-09-29T15:10:32.219

ST1

Use Case 3 : Email/Document Triage for Mailbox (2/4)

HEXAWARE

Assumptions & Dependency	Instabase APIs are accessible and support programmatic submission routing Email systems (Outlook) allow integration via API/connectors for real-time monitoring Historical labeled data exists or can be created to train classification ML models Underwriters are available to handle exception cases flagged by the AI
Input Required	Incoming broker emails with submission documents (PDFs, Word docs, Excel files) Email metadata (subject lines, sender info, body text)
Target/ Future State	**SEGMENT 1: LAUNCH CONFIGURATION (Human-in-Loop)** - AI monitors individual underwriter Outlook inboxes (addresses "brokers email Fred directly" problem) - AI classifies submission type: New Business, MTA, Inquiry - AI presents recommended classification + routing destination to underwriter - Underwriter reviews and approves/modifies via one-click approval UI - System tracks approval patterns to refine confidence thresholds **SEGMENT 2: MATURE CONFIGURATION (Graduated Autonomy)** - High-confidence classifications (>95%) auto-route without approval - Medium-confidence (80-95%) continue approval requirement - Low-confidence (<80%) flagged for manual review with explanation **DOWNSTREAM INTEGRATION:** - Auto-routed submissions trigger Instabase AI Hub ingestion for document extraction - Extends current Commercial Combined deployment to broader LOB coverage over time **REMAINING MANUAL OVERSIGHT:** - Exception handling: Underwriter validates flagged low-confidence submissions - Model retraining based on approval/override patterns
Technology & Data Considerations	Email integration layer (Microsoft Graph API for Outlook or Exchange Web Services) ML/NLP classification model for document/email categorization (could leverage pretrained ML models, Instabase ML capabilities, Azure Cognitive Services or external models) Confidence scoring mechanism with configurable thresholds for exception routing API connectors to Instabase mailboxes and workflow orchestration
Output	Classified submissions auto-routed to correct Instabase parser/workflow endpoints, confidence scores for each classification

Note: "Segments" refer to functional implementation stages for the Use Cases; "Phases" refer to TOM delivery stages (all use cases)

Slide 53

ST1 [@Anton Makarevich] Please fill this
Shreyash Tiwari, 2025-09-29T15:10:13.217

AM1 0 done
Anton Makarevich, 2025-09-30T08:12:05.382

Use Case 3: Email/Document Triage for Mailbox (3/4)

HEXAWARE

Business Value KPIs

BASELINE MTERICS:
- 5-day average quote turnaround (manual routing bottleneck)
- Brokers email individual underwriters, bypassing group mailbox
- Inconsistent triage across underwriting teams; no std routing logic
- Current: Instabase active for Comm Combined: Other LOBs manual

TARGET METRICS (12–18-month realization):
- 4-hour quote turnaround (67% reduction)
- 0.15-0.2 CR points improvement contribution
- AI classification accuracy: >86% initial, >95% mature state (Phase 2)
- Scalable process enables growth without proportional UW headcount increase

BUSINESS OUTCOMES:
- ✓ Faster broker response time
- ✓ Reduced misrouting/missed submissions
- ✓ Builds trust through HITL->graduated autonomy
- ✓ Extends to Claims email triage (reusable capability)

Risk & Mitigation

- ➢ Data exposure: Start with lowest-risk mailboxes + HIL; redact PII in prompts.
- ➢ Routing errors: Shadow mode + sampling reviews before enabling default-on suggestions.
- ➢ Adoption: Simple Outlook add-in/side-pane UX with one-click “Approve & file”.

Implementation Timelines with Standard roles requirements

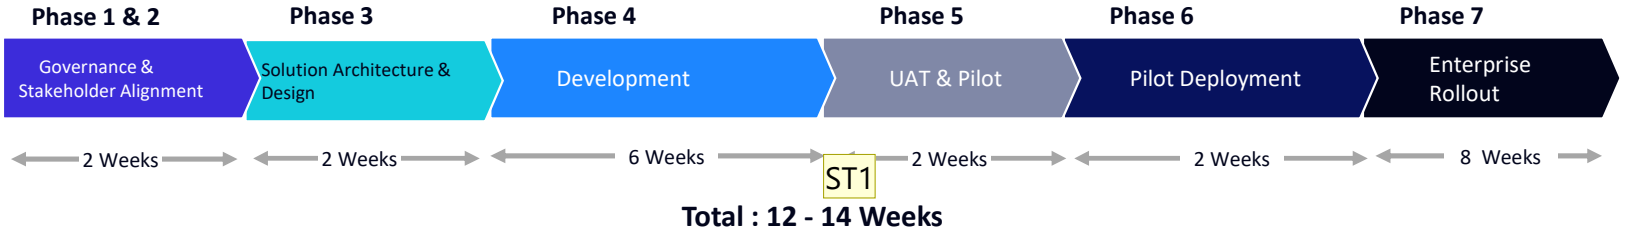

Total : 12 - 14 Weeks

Roles:

Data scientist/ML engineer (1), Backend developer (2), DevOps engineer (1), Business Analyst (1), Solution/Technical Architect (1), QA engineer (1)

High Level Cost Estimation

T-Shirt Size: SMALL-MEDIUM
(Low-code Power Platform, standard Microsoft integrations, 3 sprints, but human-in-the-loop approval UI adds design complexity)
Cost Breakdown by Category:
- Build Cost: £180-220K
- Licensing Cost: £15-25K
- Annual Run Cost: £20-30K
- Infrastructure Cost: £30-40K

Total Q1 2026 Investment: ~£265K
Cost Breakdown by Phase:
- Phases 1-3 (Governance): £20K (stakeholder validation workshops, DPIA if required)
- Phase 4 (Architecture): £30K (Power Automate design, approval UI wireframes, EA validation)
- Phase 5 (Development - 3 sprints): £120-150K (NLP model, approval interface, Instabase integration)
- Phase 6 (UAT/Pilot): £40K (10 underwriters, pilot monitoring, accuracy tracking)
- Phase 7 (Rollout): £55-75K (enterprise deployment, confidence threshold tuning, change management)
- Contingency (10%): £25K

Slide 54

ST1 [@Anton Makarevich] Please provide rough estimate of roles & timelines
Shreyash Tiwari, 2025-09-30T08:46:33.063

AM1 0 roles are updated, for the cost see my comment on #57. the timeline here is also very optimistic I would take 8 sprints (16 weeks) as a baseline for all the POCs, where first sprint - sprint 0 is design preparations, tools enablements, licensing arrangements, and sprint 8 is stabilization. 6 sprints for development itself
Anton Makarevich, 2025-09-30T10:07:14.069

Use Case 3 Implementation Plan: Email/Document Triage for Mailbox (4/4)

HEXAWARE

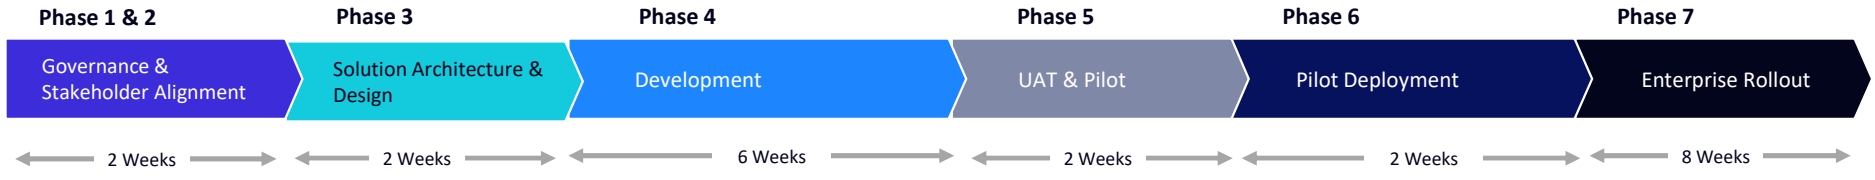

	Phase 1 & 2	Phase 3	Phase 4	Phase 5	Phase 6	Phase 7
Activities	• Underwriter engagement (10 Commercial Lines UWs) • Human-in-the-loop approval workflow design	• Power Automate flow design • Microsoft Graph API integration (Outlook) • Dynamics 365 API integration • Instabase API trigger design • Approval interface UX design	• **Sprint 1 (2 weeks):** NLP classification model training (historical emails 6-12 months) • **Sprint 2 (2 weeks):** Approval interface development + Power Automate orchestration • **Sprint 3 (2 weeks):** Instabase trigger integration + confidence scoring	• UAT with 10 Commercial Lines underwriters • Shadow mode testing (AI suggests, humans always approve) • Approval/override metrics collection • Confidence threshold calibration	• Pilot with Commercial Lines underwriters • Approval loop operational (Phase 1 launch configuration) • Performance monitoring • Confidence threshold tuning	• Rollout to all UW + Claims teams • Confidence threshold tuning (aim for >95% for auto-routing) • Gradual approval loop removal (Phase 2: high-confidence auto-routing)
Key Deliverables	• 10 underwriters identified for UAT • Approval UI design validated	• Integration architecture validated by EA • API endpoints documented • One-click approval UI wireframes • Confidence scoring framework	• Email classification model (new business, MTA, inquiry) • One-click approval interface • Outlook to Dynamics routing workflow • Instabase ingestion trigger	• 10 underwriters trained • Approval pattern analysis • Model accuracy metrics • Confidence threshold recommendations	• Pilot metrics dashboard • 5 strategic brokers monitored • Executive alerts operational	• Enterprise-wide deployment • >95% confidence auto-routing enabled • Approval loop removed for high-confidence classifications
Critical Success Factors	• Confirmed as Use Case by Graham Wilde (accelerator) • Strong underwriter buy-in for pilot	• Microsoft Graph API access confirmed • Dynamics CRM group mailbox configured • Instabase API documentation available	• Labeled training data (6-12 months historical emails) • Outlook plugin deployment permissions • Dynamics group mailbox routing rules	• Underwriter feedback on UI usability • Approval rate tracking (target >90% acceptance)	• Underwriter trust-building • Low friction approval process	• Change management across all teams • Continuous model retraining • Exception handling process mature

AM2

Use Case 4 : Instabase ingestion validation & external data enrichment(1/5)

HEXAWARE

Problem Statement	Automating validation and enrichment of data ingested through **instabase** platform.
Proposed by	Andrew Hall (CUO – Head of Technical Delivery); supported by underwriters (validation burden)
Current State	**Data Ingestion:** • Currently Instabase ingests broker slips and other relevant documents , mail attachments to capture and convert unstructured data into structured format but there is no automated validation or enrichment layer to verify accuracy before downstream systems consume the data. • Current Limitation: Instabase deployment is limited to Commercial Combined LOB. No systematic validation or enrichment workflow exists for other lines of business. **Data Validation:** • After Instabase converts broker slips to structured data, Underwriters must still perform manual external research (news, company events, asset incidents) to uncover risk signals, creating inconsistent risk assessment and potential for missed red flags. **Data Enrichment** • Ops must fix issues late in the process when validation gaps are discovered and bring them into standard format manually. No system-embedded workflow → regulatory exposure. • Copilot 365 is piloted but separate.

Use Case / Automation Potential	**Two-Segment AI Automation Approach:**
	• **SEGMENT 1: POST-INSTABASE VALIDATION (Q1 Priority)**
	  • AI-driven validation of Instabase extraction accuracy (cross-check extracted data against source slip and business rules)
	  • Automated data quality scoring with confidence metrics
	  • Exception routing for low-confidence extractions requiring manual review
	  • Scope: Commercial Combined LOB only (current Instabase deployment)
	• **SEGMENT 2: EXTERNAL DATA ENRICHMENT (Q2/3 Post-Phase 1 stabilization)**
	  • Automated risk-enrichment service queries external data sources (public news/press, corporate registries, industry feeds, sanctions/watchlists)
	  • AI flags relevant risk indicators (incidents, adverse filings, regulatory actions) with confidence scoring and source citations
	  • Enrichment results embedded in PowerUp workbench with links to sources
	  • Note: Sanctions/watchlists enforcement via dedicated compliance workflow (separate from enrichment process)"

Systems Impacted	• Instabase (trigger), PowerUp (display), external data/APIs for news/press/registry feeds; optional touchpoints S360, Copilot 365.

Note: "Segments" refer to functional implementation stages for the Use Cases; "Phases" refer to TOM delivery stages (all use cases)

Slide 56

ST1 [@Shruti Lawani] Please help us to fill this
Shreyash Tiwari, 2025-09-29T15:10:57.177

AM2 [@Shruti Lawani] @Dhanabal Subramaniam Should we mention EDP here, as it doesn't exist in "current" state?
Anton Makarevich, 2025-09-30T08:16:16.613

SL2 0 it was mentioned in the workshop by the user that it has been identified as target state. But we can reconsider whether or not to mention it here. I would go by your suggestion.
Shruti Lawani, 2025-09-30T08:20:46.242

AM2 1 yes, that's what i mean, that it probably should go to the target state, and in the current we should be very clear that there is no EDP/EDP is missing
Anton Makarevich, 2025-09-30T08:25:45.191

AM2 2 [@Shruti Lawani] what use case is it (in our master xls)?
Anton Makarevich, 2025-09-30T08:39:43.680

SL2 3 It should be 6 and 8. making the relevant changes now
Shruti Lawani, 2025-09-30T12:03:53.160

ST1

Use Case 4 : Instabase ingestion validation & external data enrichment(2/5)

HEXAWARE

Assumptions & Dependency	Instabase extraction quality is sufficient for AI validation to work effectively Broker slips follow standardized or semi-standardized formats that can be consistently parsed Reliable access to external APIs (sanctions lists, watchlists, news feeds, public data sources) is provided Predefined business rules and validation logic must be documented and agreed upon by underwriting teams
Input Required	Broker documents and their Instabase-extracted structured data (JSON/XML format); business validation rules, thresholds, and confidence score criteria; API credentials and access to external data sources (sanctions, watchlists, news).
Target/ Future State	**SEGMENT 1: POST-INSTABASE VALIDATION (Q1 2026)** • AI validates Instabase extractions against business rules with confidence scoring • Exception routing: <80% confidence flagged for manual review • Scope: Commercial Combined LOB; human oversight for exceptions **SEGMENT 2: EXTERNAL DATA ENRICHMENT (Q2-Q3 2026)** • Automated querying: Sanctions/watchlists, news/press, corporate registries • Risk indicators flagged with confidence scores and source citations • Embedded in PowerUp workbench for underwriter access **INTEGRATION:** • Instabase (trigger) → PowerUp (display) → External APIs (data) • Future EDP integration (2026-2027; not Phase 1 dependency) **EXPANSION:** Segment 1 foundation enables rollout to all UW & Claims LOBs beyond Commercial Combined **MANUAL OVERSIGHT:** Underwriter review of flagged exceptions, business rule validation
Technology & Data Considerations	AI/ML model for data validation and confidence scoring Integration layer to connect Instabase, PowerUp, S360, and external APIs Orchestration platform for exception routing and workflow management Data security and compliance controls for handling sensitive information and third-party data. Connection to Enterprise Data Platform (EDP) is part of RSA's 2026-2027 architecture roadmap but should NOT be considered a dependency for Phase 1 validation implementation.
Output	Validated and enriched broker slip data with confidence scores, flagged exceptions routed for human review

Slide 57

ST1 [@Anton Makarevich] , [@Dhanabal Subramaniam] Please help us to fill this
Shreyash Tiwari, 2025-09-29T15:11:22.378

AM1 0 done
Anton Makarevich, 2025-09-30T09:46:07.597

Use Case 4: Instabase ingestion validation & external data enrichment(3/5)

HEXAWARE

Business Value KPIs

SEGMENT 1 (Q1 2026 - Validation Engine):

Baseline Metrics:
- 100% manual validation of Instabase extractions
- 15-20 min/submission verification time
- No confidence scoring for extracted data

Target Metrics (6-9 months realization):
- 80-90% auto-validation with confidence scores
- Only <80% confidence requires human review
- 10-15 min/submission time savings (reusable layer)

SEGMENT 2 (Q2-Q3 2026 - Enrichment Layer):

Baseline Metrics:
- 50% manual sanctions check coverage (inconsistent)
- Adhoc news/ risk research (no systematic workflow)
- No enrichment embedded in PowerUp workbench

Target Metrics (Q2-Q3 26) :
- 100% automated sanctions/watchlist checks
- Relevant risk indicators in PowerUp with citations
- Embedded enrichment & compliance workflow

BUSINESS OUTCOMES:
- 100% automated sanctions coverage
- Enhanced UW quality
- Scalable foundation for all LOB
- Risk based human review

Risk & Mitigation

- Risk: False positives in validation/sanctions → incorrect decisions
- Mitigation: Human-in-the-loop for low-confidence; feedback loop for model improvement

Implementation Timelines with Standard roles requirements

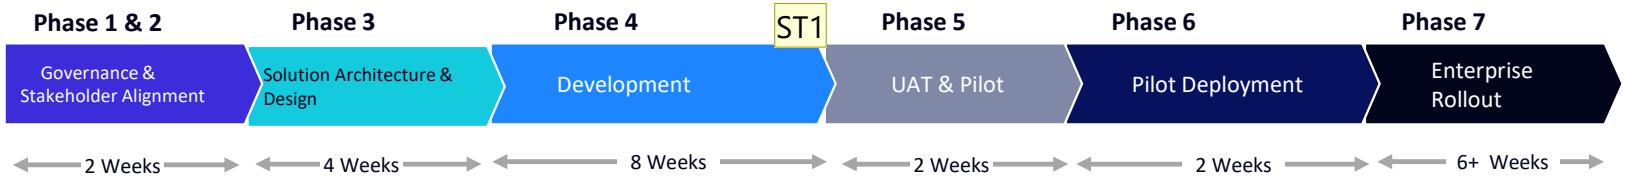

Total Duration: 16-18 weeks

Segment 1 (Weeks 1-12): Post-Instabase Validation Engine

Segment 2 (Weeks 13-18): External Data Enrichment Layer

Roles:

Data scientist/ML engineer (1), Backend developer (2), DevOps engineer (1), Business Analyst (1), Solution/Technical Architect (1), QA engineer (1)

High Level Cost Estimation

T-Shirt: LARGE (Two-Segment)

Segment 1 (Q1 2026): £625-775K ✅
- Build Cost: £500-600K
- Licensing: £0 (existing Azure/Instabase)
- Annual Run: £50-75K
- Infrastructure: £75-100K

Segment 2 (Q2-Q3 2026): £725K-1.0M ⏳
- Build Cost: £400-500K
- Licensing (ONE-TIME): £200-300K 🚨
- Annual Run: £75-125K 🚨
- Infrastructure: £50-75K

Note: "Segments" refer to functional implementation stages for the Use Cases; "Phases" refer to TOM delivery stages (all use cases)

Slide 58

ST1 [@Anton Makarevich] Please provide rough estimate of roles & timelines
Shreyash Tiwari, 2025-09-30T08:46:51.531

AM1 0 [@Shreyash Tiwari] the only thing we can do here is to "estimate" build based on the t-shirt sizing and using a generic formula for all the use cases (basically what the portal does currently), in this scenario we will have some justification at least. everything else will be a very wild guess (based on the level we are currently at)
Anton Makarevich, 2025-09-30T09:48:26.092

AM1 1 roles have been updated, for the timeline, again at least 8 sprints (4 months) with selected integrations
Anton Makarevich, 2025-09-30T11:17:14.695

Use Case 4 Implementation Plan: Instabase ingestion validation & external data enrichment (4/5)

DP1

Iteration 1 (Weeks 1-12): Post-Instabase Validation Engine

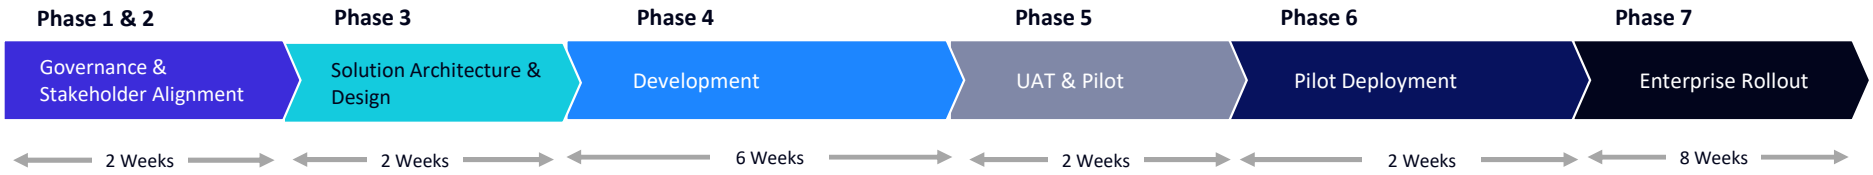

	Phase 1 & 2	Phase 3	Phase 4	Phase 5	Phase 6	Phase 7
Activities	• Andrew Hall (CUO) validation: "Critical gap after Instabase conversion" Underwriter engagement on validation rules • Business rules workshop	• Instabase output schema analysis • Validation rule framework design • Data quality scoring model architecture • Exception routing workflow design & PowerUp workbench integration	• **Sprint 1 (2 weeks):** Validation engine core + Instabase API integration • **Sprint 2 (2 weeks):** Data quality scoring model + confidence metrics • **Sprint 3 (2 weeks**): Exception handling + PowerUp workbench display	• UAT with 5 Commercial Combined underwriters • Validation rule accuracy testing • Exception workflow testing • Confidence threshold calibration	• Pilot with Commercial Lines underwriters • Approval loop operational (Phase 1 launch configuration) • Performance monitoring • Confidence threshold tuning	• Rollout to all UW + Claims teams • Confidence threshold tuning (aim for >95% for auto-routing) • Gradual approval loop removal (Phase 2: high-confidence auto-routing)
Key Deliverables	• Andrew Hall executive sponsorship • Validation rules catalogue (documented & agreed) • Confidence scoring criteria	• Instabase JSON/XML schema documented • Validation rule engine design • Exception handling workflow • PowerUp integration specifications	• Email classification model (new business, MTA, inquiry) • One-click approval interface • Outlook to Dynamics routing workflow • Instabase ingestion trigger	• 10 underwriters trained • Approval pattern analysis • Model accuracy metrics • Confidence threshold recommendations	• Pilot metrics dashboard • 5 strategic brokers monitored • Executive alerts operational	• Enterprise-wide deployment • >95% confidence auto-routing enabled • Approval loop removed for high-confidence classifications
Critical Success Factors	• Confirmed as priority gap by CUO (accelerator) • Strong underwriter participation	• Instabase schema well-documented • Business rules clearly defined • DO NOT WAIT for EDP integration (critical to avoid roadmap dependency)	• Instabase API access and documentation • PowerUp workbench enhancement permissions • Business rule translation to code	• Underwriter feedback on UI usability • Approval rate tracking (target >90% acceptance)	• Underwriter trust-building • Low friction approval process	• Change management across all teams • Continuous model retraining • Exception handling process mature

Slide 59

DP1 [@Anton Makarevich] and [@Mukundan Madhavan]
Donna Pahel, 2025-10-08T17:14:01.145

Use Case 4 Implementation Plan: Instabase ingestion validation & external data enrichment (5/5)

HEXAWARE

Iteration 2 (Weeks 13-18): External Data Enrichment Layer

	Phase 1 & 2	Phase 3	Phase 4	Phase 5	Phase 6	Phase 7
Activities	• Underwriter engagement (10 Commercial Lines UWs) • Human-in-the-loop approval workflow design	• Power Automate flow design • Microsoft Graph API integration (Outlook) • Dynamics 365 API integration • Instabase API trigger design • Approval interface UX design	• **Sprint 1 (2 weeks):** NLP classification model training (historical emails 6-12 months) • **Sprint 2 (2 weeks):** Approval interface development + Power Automate orchestration • **Sprint 3 (2 weeks):** Instabase trigger integration + confidence scoring	• UAT with 10 Commercial Lines underwriters • Shadow mode testing (AI suggests, humans always approve) • Approval/override metrics collection • Confidence threshold calibration	• Pilot with Commercial Lines underwriters • Approval loop operational (Phase 1 launch configuration) • Performance monitoring • Confidence threshold tuning	• Rollout to all UW + Claims teams • Confidence threshold tuning (aim for >95% for auto-routing) • Gradual approval loop removal (Phase 2: high-confidence auto-routing)
Key Deliverables	• 10 underwriters identified for UAT • Approval UI design validated	• Integration architecture validated by EA • API endpoints documented • One-click approval UI wireframes • Confidence scoring framework	• Email classification model (new business, MTA, inquiry) • One-click approval interface • Outlook to Dynamics routing workflow • Instabase ingestion trigger	• 10 underwriters trained • Approval pattern analysis • Model accuracy metrics • Confidence threshold recommendations	• Pilot metrics dashboard • 5 strategic brokers monitored • Executive alerts operational	• Enterprise-wide deployment • >95% confidence auto-routing enabled • Approval loop removed for high-confidence classifications
Critical Success Factors	• Confirmed as Use Case by Graham Wilde (accelerator) • Strong underwriter buy-in for pilot	• Microsoft Graph API access confirmed • Dynamics CRM group mailbox configured • Instabase API documentation available	• Labeled training data (6-12 months historical emails) • Outlook plugin deployment permissions • Dynamics group mailbox routing rules	• Underwriter feedback on UI usability • Approval rate tracking (target >90% acceptance)	• Underwriter trust-building • Low friction approval process	• Change management across all teams • Continuous model retraining • Exception handling process mature

HEXAWARE

Illustrative Future-State Use Case: Agentic AI-Enabled Commercial Property Underwriting

This forward-looking example illustrates how RSA's maturing AI capabilities could evolve toward fully agentic workflows — connecting insights, decisions, and compliance in real time across the underwriting value chain.

SUBMISSION	EXTRACTION	ENRICHMENT	ASSESSMENT			UNDERWRITING PORTAL	
Intelligent Email Screening	**Intelligent Data Ingestion**	**Data Enrichment**	**Comparative Analytics**	**Triage: Appetite**	**Triage: Propensity**	**Underwriting (Decisioning & Pricing)**	**Insurance Core System Integration**
• Orchestrates channel monitoring • Summarizes broker email • Classifies submission vs inquiry	• Autonomous Parsing Agent • Drafts missing data request • Field-level confidence scoring, anomaly detection	• Enrichment Agent merges geo, peril, sanctions results • Creates Evidence Book narrative • Screens hazards, Risk classifications	• MPL/PML calculation flow • Visual risk profile summary • Fraud detection & Loss Ratio Prediction models	• Planner Agent decides alternative coverage path • Summary of Appetite decision • Appetite classification model	• Propensity to Bind Model Agent • Quote prioritization explainer • Expected premium vs. effort model	• UW Co Pilot with Pricing & Wording Agent • Quote Pack Generator • Rate adequacy classification	• Automates PolicyCenter updates • Personalised bind/decline • Final compliance checks

SUCCESS METRICS / KPIs

Ingestion success %	Field-level confidence %	Enrichment coverage %	Peer-set size (comparables)	Auto-triage precision vs UW	Propensity model lift	Quote TAT	STP %

BUSINESS OUTCOMES

- **Faster turnaround (minutes → hours)**
- **Higher accuracy, fewer errors**
- **Improved win rates & hit ratios**
- **Scalable underwriting capacity**

GOVERNANCE & TRUST

- **Evidence Book: source + confidence tracking**
- **Guardrails: sanctions, fairness, PII/PHI**
- **Decision Ledger: full AI + UW audit trail**
- **Feedback Loop: continuous model learning**

Agentic AI | Generative AI | Discriminative AI

***Note**: Illustrative outside-in example for reference; not part of RSA's current delivery scope.*

HEXAWARE

Architecture & Data Enablement

AM1

Architecture & Data Enablement

Building the Intelligent Core for Scalable AI

HEXAWARE

RSA's architecture and data transformation is already underway — establishing the governed and modular foundation that will power enterprise-wide AI adoption.

AM2

- **Enterprise Data Platform** – Establishes a consistent, governed data foundation to eliminate silos and enable trusted, reusable datasets.
- **Architecture Refactoring** – Transforms a fragmented "swamp" into a cleaner, modular "bubble" model built for maintainability and scale.

Building on RSA's existing momentum, two additional capabilities can enhance scalability and fault-tolerance of the architecture and help with the timely delivery of new features and use cases:

- **Event-Driven Design** – Prepares RSA for scalable, flexible, AI-driven workflows through asynchronous, real-time orchestration.
- **Hexaware Partnership** – Builds on RSA's existing momentum to operationalize AI safely and efficiently across the enterprise.

AM3

This foundation enables RSA to move from isolated pilots to a unified, intelligent ecosystem where agentic AI can act, learn, and scale responsibly.

Slide 63

AM1 [@Donna Pahel] please validate if this text aligns with your thoughts
Anton Makarevich, 2025-10-02T09:06:04.245

AM2 2 fiirs items are "underway", the third and fourth is what we recommend additionally. i think that should be clearly distinguished
Anton Makarevich, 2025-10-08T13:32:43.959

AM3 added this to address the comment above [@Donna Pahel]
Anton Makarevich, 2025-10-08T13:42:00.197

HEXAWARE

Core Components of Agentic AI System

Agentic AI systems integrate orchestration, data, models, and governance layers — enabling RSA to build, deploy, and manage intelligent agents that reason, act, and improve autonomously.

Orchestration

Workflow and policy engine coordinating perception-plan-act loops and multi-agent hand-offs.

Agentic Framework

Abstractions simplifying development with standardised interfaces for tools, memory, planning, routing logic, and safety mechanisms.

AI Model

Cognitive core (including an LLM, SLM or a predictive model) as reasoning engine for semantic understanding.

Compute & Runtime

Infrastructure for executing workloads efficiently: serverless functions, containers, and batch workers designed to handle variable and bursty AI workloads.

Data Plane

Data stores for external knowledge bases, and feature stores for structured and unstructured data.

Integrations

Bridge connecting to external systems: business systems, data sources, event streams, and APIs enabling agents to perform actions and retrieve real-time information.

Safety

Ensuring responsible operation through content filters, policy enforcement, and safeguards to prevent hallucination, data leakage, and harmful behaviours.

Governance

Oversight layer that defines audit trails, access controls, and compliance measures to ensure transparency, accountability, and responsible lifecycle management of agents and models.

Observability

Insights for monitoring and debugging: distributed tracing, telemetry data collection, cost/latency monitoring, and experiment tracking for continuous improvement.

HEXAWARE

Orchestration of Multi-Agent Systems

Managing coordination across intelligent agents requires balancing centralized control with dynamic autonomy — blending orchestration and event-driven choreography for resilience and scalability.

Centralised Approach (Orchestration)	Event-Driven Approach (Choreography)
Agents communicate directly with each other (A2A). A central workflow orchestration system such as Apache Airflow, AWS Step Functions, Azure Logic Apps or an orchestrator integrated with the agentic framework is required to coordinate interactions.	Agents publish events and listen for matched events. Agents are independent from each other and can adjust their strategies dynamically. Implements choreography pattern rather than orchestration.

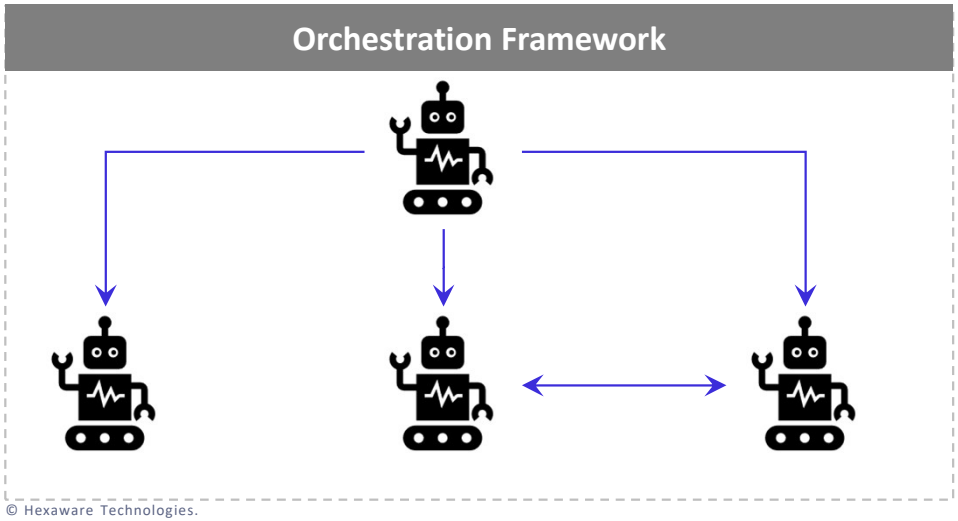

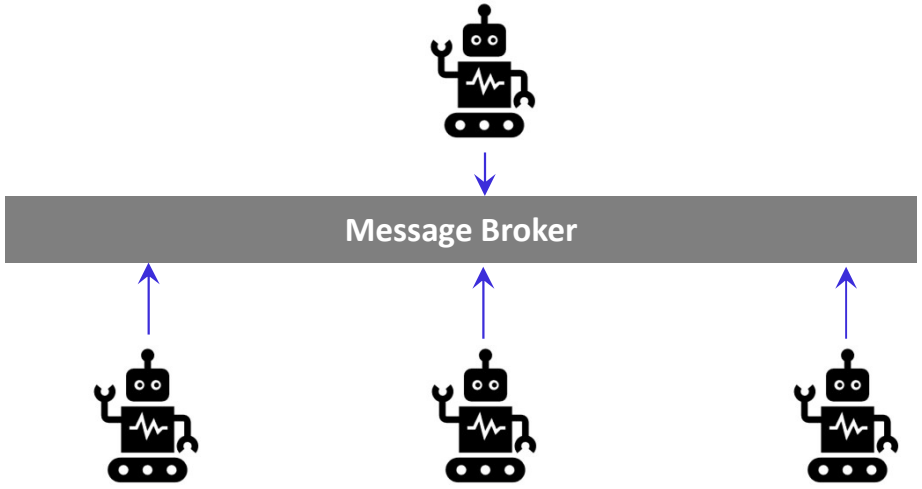

HEXAWARE

Choosing the Orchestration Approach

RSA's environment benefits from a **hybrid model** — retaining central oversight for critical workflows while embracing event-driven flexibility for real-time, distributed intelligence.

Centralised Orchestration	Event-Driven Orchestration
• Strong compliance and audit requirements • Complex business logic with strict sequencing is needed • Straightforward debugging and troubleshooting • Team has limited distributed systems experience • Workflow visualization for business stakeholders is critical	• High scalability and fault tolerance are required • Agents need autonomous decision-making capabilities • System must handle unpredictable workload patterns • Different agents may be owned by different team • Real-time responsiveness is critical

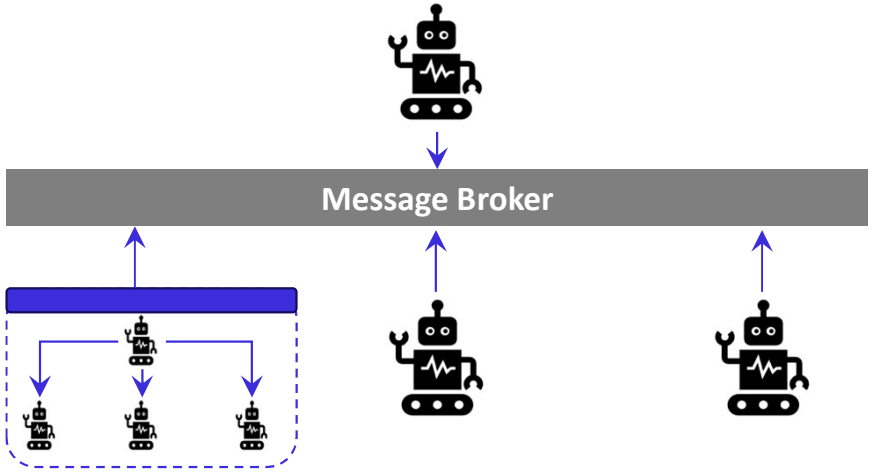

Recommended Option

Hybrid Approach

Many production systems may combine both patterns, using centralized orchestration for critical business workflows while relying on global event-driven architecture

Agentic Frameworks

Agentic frameworks provide the backbone for how agents think and operate — linking prompts, models, tools, and planning logic into cohesive, reusable architectures.

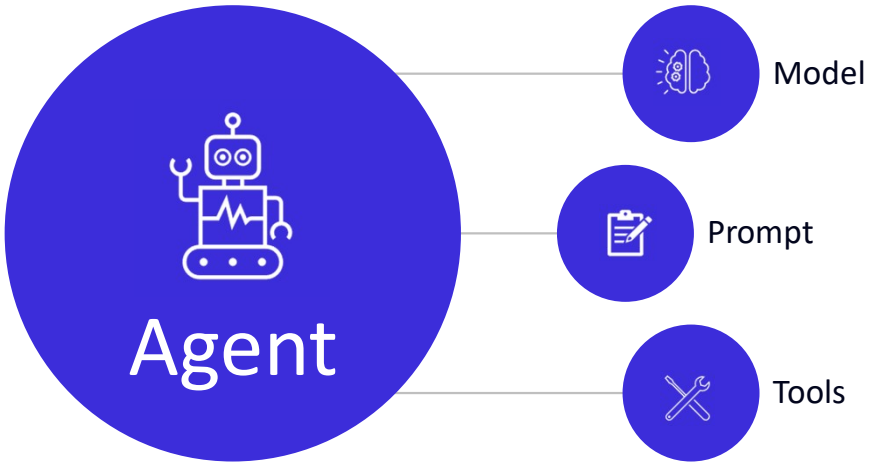

Common Components

- Agent Management: Lifecycle management, state tracking and agent registration
- Model Integration: Standardised interfaces for different AI models, routers support
- Prompt Assembly: Dynamic prompt construction and template management
- Tool Integration: Unified interface for external tool access
- Memory Management: Short-term and long-term memory systems
- Planning Logic: Task decomposition and execution planning
- Routing Logic: Decision-making for task delegation and agent selection
- Safety Mechanisms: Built-in guardrails and validation systems

AM1

Framework Selection Considerations

HEXAWARE

By aligning with Azure's Semantic Kernel while allowing for LangChain and CrewAI experimentation, RSA can accelerate delivery without locking itself into proprietary or custom frameworks.

Open-Source Frameworks (just a few mentioned)	
Semantic Kernel	Microsoft-backed, ready for enterprise-grade solutions, plugins-based, multi-language SDKs
Lang* Family	Very mature, production proofed, comprehensive ecosystem, extensive tool integrations, rapid iteration
CrewAI	Easy to start with, multi-agent focus, role-based agent design, structured workflows
Custom Framework Development	
Pros	Full control, optimized for specific use cases
Cons	High development cost, ongoing maintenance burden, limited ecosystem
Commercial Platforms	
Pros	Professional support, enterprise features, compliance certifications
Cons	Vendor lock-in, licensing costs, less customization flexibility

Semantic Kernel

crewai

LangChain

Recommended Option

- **Semantic Kernel:** the best to be used within Microsoft/Azure ecosystem
- **LangChain/Graph**: for AWS-based and cloud agnostic solutions
- **CrewAI**: for quick POCs and experiments

- Do not invest into custom frameworks unless necessary

Slide 68

AM1 [@Donna Pahel] do you want to remove that? I'm fine, as we now have this in the description
Anton Makarevich, 2025-10-08T13:46:42.189

Compute Services

HEXAWARE

Flexible compute choices — from containerized Kubernetes clusters to low-code and serverless functions — allow RSA to scale AI workloads efficiently while preserving architectural consistency.

Cloud Infrastructure				
Serverless Functions	Managed App Services	Container Apps	Kubernetes	Virtual Machines
• AWS Lambda • Azure Functions	• Azure App Service • AWS Elastic Beanstalk	• Azure Container Instances • AWS App Runner • AWS ECS (Fargate)	• Azure Kubernetes Service • AWS Elastic Kubernetes Service • Managed Kubernetes providers	• Azure Virtual Machimes • AWS EC2
-		**Possible with On-Prem Infrastructure**		

Recommended Option

Container orchestration with Kubernetes: Highly scalable and cloud agnostic in its core

Serverless (i.e. Azure Functions) and **low code** (i.e. Logic Apps, Power Platform) for specific use cases

Integrations with Tools and External Services

HEXAWARE

Agentic ecosystems thrive on interoperability — standardized APIs, message brokers, and protocols like MCP and A2A enable seamless communication between agents and systems.

Example of Integrations

API Endpoints
Direct HTTP/REST API calls to external services
REST APIs, GraphQL, Webhooks

Agent to Agent
Task delegation to other agents
Sync over HTTP or Async over events mesh

Database Connectors
Direct database access for data retrieval and manipulation
SQL, NoSQL, Vector DBs

File System Access
Read/write operations on local or cloud file systems
Local files, S3, Google Drive

Example of Integration Protocols

Model Context Protocol (MCP)
Standardised protocol for secure context sharing between AI models and tools
Anthropic MCP, local/remote servers

Agent to Agent Protocol (A2A)
Standardised protocol for secure synchronous communication between AI agents
•Google A2A, agent servers

OpenAPI Specification
Standard for describing REST APIs, LLM integration requires description of usages scenarios
OAS, Swagger, API documentation

AsyncAPI Specification
Asynchronous communication through message brokers
RabbitMQ, Apache Kafka, Redis

Recommended Option

Always follow established open standards and protocols

AI Model Types: Strengths and Limitations

HEXAWARE

Each model type—Predictive ML, Large Language Models (LLMs), and Small Language Models (SLMs)—offers distinct advantages and trade-offs; understanding these helps RSA align the right model with the right use case.

Predictive (ML) Model	Large Language Model	Small Language Model
Pros:		
• Highly optimized for specific tasks • Fast inference and training • Low computational requirements • Interpretable and explainable • Proven reliability in production • Well-understood behaviour • Cost-effective at scale	• Exceptional natural language understanding and generation • Strong reasoning capabilities across diverse domains • Can handle complex, multi-step tasks • Rich contextual understanding • Code generation and debugging capabilities • Emergent abilities at scale	• More efficient resource usage • Faster inference than LLMs • Easier to fine-tune and customize • Better privacy (can run locally) • Lower operational costs • More predictable behaviour • Suitable for edge deployment
Cons:		
• Narrow task specialization • Requires extensive feature engineering • Poor to no natural language capabilities • No reasoning or planning abilities • Rigid input/output formats • Needs separate models for each task • Tools integration capabilities are limited	• Extremely high computational costs • Slow inference times • High memory requirements • Potential for hallucinations • Difficult to fine-tune efficiently • Black box interpretability • May be overkill for simple tasks	• Limited reasoning capabilities • Reduced knowledge breadth • Less sophisticated language understanding • Weaker few-shot learning • May struggle with complex tasks • Limited emergent abilities • Narrower application scope

AI Models: Capabilities and Cost

HEXAWARE

Balancing predictive ML, large language models, and small language models ensures RSA can match intelligence depth with operational efficiency across diverse business contexts.

	Predictive (ML) Model	Large Language Model	Small Language Model
Model Size	Vary a lot, but limited	Up to trillions of parameters	Millions to several Billions of params
Training Cost	Low to Moderate	Very High	Moderate to High
Hardware Requirements	CPUs sufficient	High-end GPUs	Mid-range GPUs
Energy Consumption	Low	Very Hight	Medium
Use When	• Task is well-defined and narrow • Performance and efficiency are paramount • Interpretability is required • Working with structured data primarily	• Natural language interaction is critical • Complex reasoning is required • Flexibility across domains is needed • Budget allows for high computational costs	• Need balance between capability and efficiency • Deploying to edge devices or mobile • Privacy concerns require local processing • Task complexity is moderate

Recommended Option

Start with available **LLMs**, prioritise open-sourced and privately hosted models
Move towards finetuned and **specialised smaller models** when available
Traditional **ML** models trained on organisational data are very relevant in many real-life scenarios

Model Routing Systems

HEXAWARE

Model routers intelligently allocate tasks to the most suitable model based on cost, complexity, and context — maximizing performance while managing efficiency.

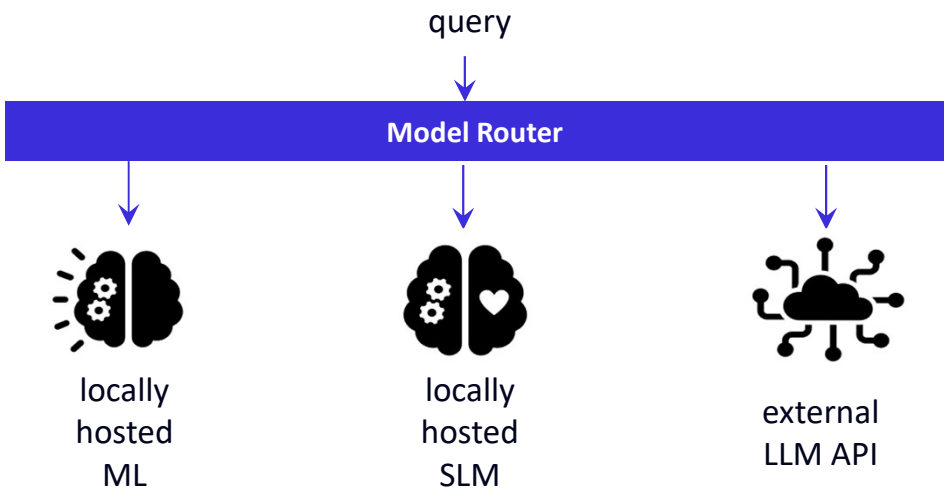

How They Work

1.Input analysis – The router inspects the request (length, language, domain, complexity, intent).

2.Model selection – It decides which model to use (fast but cheap, or slower but more accurate).

3.Execution – The chosen model generates the response.

4.(Optional) Fallback/Ensemble – If the first model fails or seems low-quality, another model may be tried.

Recommended Option

- Prefer routing at the agent level, making agents more specific by using dedicated models per agent
- Allow access to multiple models only for agents where it's necessary in complex scenarios

(such as orchestrator or coordinator agents)

Safety and Governance

Built-in safety, policy, and governance mechanisms form the ethical and operational backbone of RSA's AI architecture — ensuring transparency, fairness, and compliance at every step.

Content Safety	Data Privacy and Security
• **Input Filtering:** Malicious prompt detection and sanitization • **Output Validation:** Harmful content detection and blocking • **Bias Detection:** Algorithmic fairness monitoring and correction • **Toxicity Screening:** Real-time content toxicity assessment	• **PII Detection and Masking:** Automatic identification and protection of sensitive data • **Encryption:** End-to-end encryption for data in transit and at rest • **Access Control:** Role-based access control (RBAC) and fine-grained permissions • **Data Residency:** Geographic data storage compliance
Policy Enforcement	**Model Governance**
• **Business Rules Engine:** Automated policy compliance checking • **Approval Workflows:** Human-in-the-loop validation for decisions • **Rate Limiting:** Prevent abuse and resource exhaustion • **Audit Trails:** Comprehensive logging of all agent actions and decisions	• **Model Versioning:** Systematic model lifecycle management • **A/B Testing:** Controlled model performance comparison • **Drift Detection:** Monitoring for model performance degradation • **Explainability:** Interpretable AI for decision transparency

Observability

HEXAWARE

Observability transforms AI from a black box into a measurable system — tracking latency, cost, accuracy, and ROI to sustain trust and accountability at scale.

Performance Monitoring	Distributed Tracing
• **Latency Tracking:** End-to-end response time measurement • **Throughput Metrics:** Request processing rates and capacity utilization • **Error Rates:** Failure detection and categorization • **Resource Utilization:** CPU, memory, and GPU usage monitoring	• **Request Tracing:** Full request tracking across services • **Agent Interaction Mapping:** Visualization of multi-agent communication patterns • **Dependency Analysis:** Understanding of service interdependencies • **Performance Bottleneck Identification:** Automated performance issue detection
Business Metrics & Analytics	**Alerting and Incident Management**
• **Task Completion Rates**: Success metrics for agent objectives • **Cost Analytics:** Detailed cost breakdown and optimization opportunities • **ROI Tracking:** Business value measurement and reporting • **Usage Analytics**: User behaviour and system usage patterns	• **Proactive Alerting:** Threshold-based and anomaly detection alerts • **Incident** Response: Automated escalation and notification workflows • **Root Cause Analysis:** Automated investigation and diagnosis tools • **Recovery Procedures:** Automated remediation and failover mechanisms

Additional Infrastructure Considerations

HEXAWARE

Core enablement tools like vector search, CI/CD, secrets management, and event meshes form the connective tissue that keeps RSA's AI systems secure, responsive, and reliable.

Secrets & Configuration

Securely managing sensitive information and system configurations for agentic AI systems.

Examples:
Azure Key Vault, AWS Secrets Manager, HashiCorp Vault

Message Broker or Events Mesh

Facilitating asynchronous interactions and managing complex workflows through event-driven architecture.

Examples:
Amazon EventBridge, Azure Service Bus, Apache Kafka

Cognitive Search

Enabling long-term memory and RAG in agentic systems by converting data into vector embeddings for semantic search.

Examples:
Azure AI Search, OpenSearch, Milvus

CI/CD

Facilitating efficient, reliable, and safe development of the AI enabled workflows

Examples:
GitHub Actions, GitLab CI, Azure DevOps, ArgoCD

Deployment Options

RSA can deploy AI workloads flexibly across Azure, AWS, or hybrid models — prioritizing consistency, portability, and data sovereignty while leveraging existing investments.

Microsoft Azure

- **Azure AI Foundry** for model management
- **Azure ML** for training and deployment
- **Azure OpenAI** for LLM integration
- **Azure Service Bus** for event-driven communication
- **Azure Kubernetes Service** for container orchestration

AWS

- **Amazon Bedrock** for foundation model access
- **Amazon Sagemaker** for training and deployment
- **Amazon EventBridge** for event orchestration
- **AWS Step Functions** for workflow coordination
- **Amazon EKS** for container orchestration

Cloud-Agnostic

- **Microservices architecture** with container orchestration
- **Multi-cloud portability** and vendor independence
- **GitOps-based** deployment workflows
- **Private model hosting** for sensitive applications and **data sovereignty**

Reference Architecture: Azure/Hybrid Solution

HEXAWARE

A reference architecture anchored in Azure demonstrates how governance, orchestration, and data layers come together to support both traditional and agentic workloads.

Governance: Security, Compliance and AI Ethics (RBAC, Policy, Responsible AI etc)

Observability: Monitoring, Logging, and Performance Analytics (Azure Monitor, Application Insights, Logs Analytics etc)

Experience API: Centralized API Gateway and Developer Experience (Azure API Management)

Message Broker: Event-Driven Communication and Orchestration (Azure Service Bus or Event Hubs/Kafka)

Microservices: Container-based Application Deployment (Azure Kubernetes Service)

- Single Agent Service
  Semantic Kernel
- Multi-Agent Workflow
  Semantic Kernel
- Traditional Microservice
  Java / .NET / Node.js etc

Integration Services

AI Models: Cloud and Private Model Infrastructure

- AI Foundry
  Model Router Customized Models
  Model Catalog & Registry
- Azure OpenAI API
  GPT-5
  GPT-4o
- Private/Edge Models
  On-Premises GPU Clusters
  Edge Computing Nodes

Data Plane: Enterprise Data Platform (Data Lake, SQL, Vector, ETL)

High-Level Technical Architecture (Azure)

HEXAWARE

RSA's Azure-based implementation blueprint brings the architecture to life — connecting AKS, Service Bus, AI Foundry, and the Enterprise Data Platform into one cohesive ecosystem.

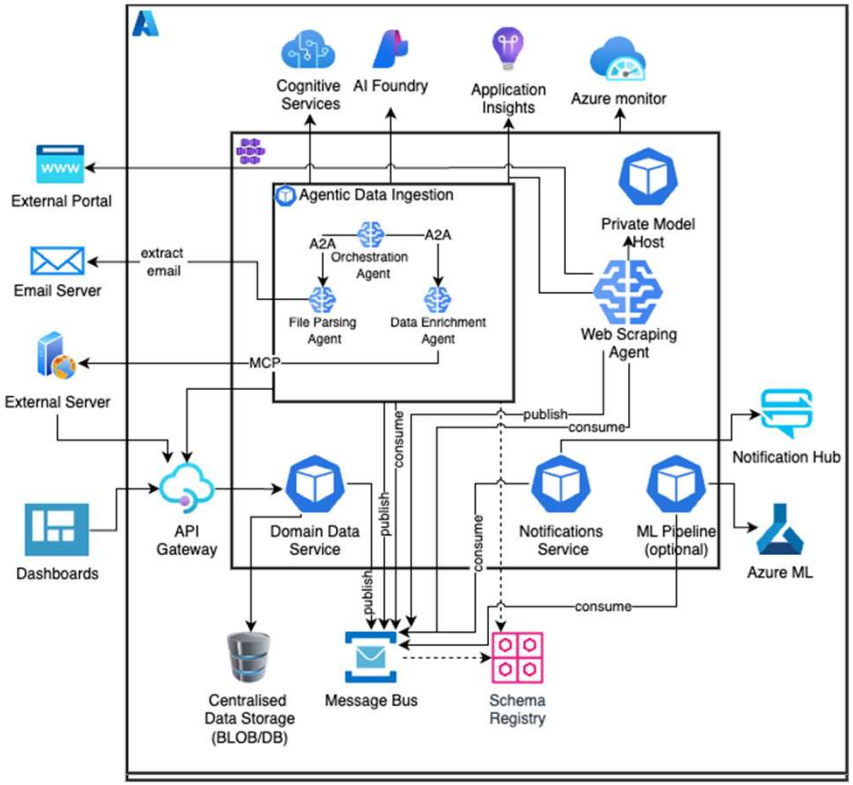

Core Infrastructure

- **Azure Kubernetes Service (AKS)** orchestrates all microservices including agents, agentic workflows, and custom models for scalable container management
- **API Management Gateway** provides unified contract for integrations
- **Azure Service Bus (Event Hubs** or **Kafka** as alternatives) enables reliable pub/sub messaging decoupling microservices
- **Centralised Storage** (Blob/DBs) serves as single source of truth for data
- **IA Foundry** to access and route AI models and APIs
- **Cognitive Service** to provide capabilities like OCR, Computer vision, Audio transcribing etc
- **Open Telemetry** compatible services for observability (either Managed or external)

HEXAWARE

Placeholder for next steps slide

Thank You

HEXAWARE

HEXAWARE

Current State Assessment

AS1

HEXAWARE

Assessing RSA's AI Readiness: Strengths, Gaps, and Priorities for Scale

RSA's AI capabilities show strong intent but remain in the early stages of operational maturity — highlighting the need for disciplined delivery, data enablement, and integrated governance.

Capability Levers

	Incubating	Innovating	Maturing	Optimizing	Transforming
Process	Lack of an organisational AI strategy with a few use cases being experimented, but no view on ROI	Early Innovating — pilots exist, ROI not visible, manual rekeying dominates Start on planning the use of AI in a subset of organisation and POC signed off	AI strategy defined along with Executive sponsorship. ROI expectations are clear	AI strategy execution, aligning with digital, data and technology strategy. ROI metrics defined	AI strategy as a DNA of organisational strategy. AI based operating model and continuous improvement
Data Management	Data Silos across different business units with no clear definition of data requirements for AI	Early Innovating — fragmented datasets, limited lineage, silos persist Clarity on AI driven data requirements with efforts to create a common data store with limited access	Core Data set required for AI solutions is in place. Strategic priority to build data infrastructure for AI strategy	Extensive, upto date and usable data for AI solutions. Strategic systems connect to a common data platform	Data platform fundamental to core business operations. Automated & Self-Service data access
Technology	No investment on an AI technology, however DevOps, RPA and Analytics exists. Lack of clarity on what is needed	Early Innovating — core tools available but not embedded in workflows Ability to provision Cloud or On-premise infrastructure for AI. Model training & deployment still happens manually	AI Deployment architecture and development tools standardised & implemented. Automation of access and resource allocation	AI model deployment increases across BU. Centralised monitoring & audit of AI models for compliance & performance management	State of the art AI infrastructure standardized and efficient. AI used to manage the technology infrastructure itself
People	No defined AI roles & responsibilities. Lack of AI literacy across the business & data teams	Mid-Innovating — AI champions emerging, literacy still low, adoption ~5% of staff Some roles & responsibilities exist but experimentation on right way to organise AI. PoC have started in isolation	Enterprise level AI roles defined. Dedicated CoE created to provide skills. Business leaders communicated on the vision for AI and focus on AI literacy	Clearly defined roles & responsibilities with well-defined KPIs. COE formalised with an organisation wide mandate. Increased AI literacy	High degree of AI literacy with AI first mind set. AI integrated in all roles including executive level. Organisation working model is AI focused
Governance	Employees & leadership being educated about responsible AI. Identification of major risks with priority use cases	Mid-Innovating — framework and council aligned, but not yet operationalized Business, technical & risk team share understanding on responsible AI. Guiding principles of AI adoption beyond legal framework	Guiding principles are converted into specific performance metrics with centralised reporting. Commitment to AI governance is formalised	Formal centralised and auditable process, policy & technology for AI practice. Risk considerations at a model level	Governance goes beyond regulatory compliance. Risk defence and trust are a competitive advantage for applying AI

Stages of Maturity

Slide 83

AS1 This table and overlapping elements should be resized to fit the slide layout
Anna Shymchenko, 2025-10-02T10:55:06.246

VA2

Current State Assessment

What We Heard From Stakeholders

Across stakeholder conversations, leaders and frontline teams emphasized the need to move beyond experimentation toward scaled, measurable AI outcomes.

Key Themes:

- **High AI demand, uneven readiness:** Enthusiasm for AI spans Claims, Underwriting, HR, and IT, but capability maturity varies widely.
- **Fragmented delivery landscape:** Multiple parallel pilots exist, but without a unified roadmap or ownership structure.
- **Data and integration bottlenecks:** Inconsistent data quality, manual reconciliation, and limited interoperability remain top barriers to scale.
- **Pressure for tangible results:** Leadership and brokers are expecting visible performance improvements and efficiency gains.
- **Skills and confidence gap:** Many employees want to use AI tools but lack guidance, governance clarity, or confidence in doing so.

> "We don't just need experiments — we need outcomes we can show the Board and our brokers."

> "Data is everywhere and nowhere — we spend more time reconciling than analyzing."

> "Our people want to use AI, but most don't know where to start or what's allowed."

> "AI feels like a set of disconnected initiatives rather than a coordinated capability."

Slide 84

VA1 4 Key messages on the Left.. And 3-4 verbatim wordings from those who cooperated the best with us.. Wrt their vision or their AI aspiration..
Vivek Arya, 2025-09-04T11:49:53.225

VA2 This slide should cover in crisp the key messages from most imp stakeholders
Vivek Arya, 2025-09-04T11:50:21.920

AS1

Current State Assessment

Rationale Behind the Maturity Ratings

HEXAWARE

RSA's AI maturity is constrained by fragmented processes, siloed data, and limited integration — AI exists in pilots and tools but is not embedded in core workflows or operating models.

Existing Eco-System

Process

AI initiatives are fragmented and lack enterprise coordination

- AI initiatives fragmented; POCs are custom-built, with little to no enterprise reuse
- Underwriters and claims teams rekey and manually route data across disconnected systems

Data

Data is siloed, inconsistent, and inaccessible for scalable AI.

- Policy data scattered across UKRIS, Instabase, legacy PAS, and delegated portals
- HR data requires ~10 hours of manual cleansing per monthly refresh due to fragmentation across SAP and other sources

Technology

The tech stack lacks orchestration, reuse, and integration.

- Tools (e.g., Copilot, CallMiner) piloted but not embedded in workflows
- Lack of orchestration prevents reuse, integration, and monitoring

People

AI literacy and role clarity are low, constrained delivery capacity

- Claims handlers read long, unstructured notes to understand context instead of acting on AI-summarized insights
- Teams rely on manual portal scraping and Excel trackers for broker pipeline generation due to lack of AI-driven opportunity surfacing

Governance

Council in place; self-attestations started; risk/compliance guardrails advancing

- There's a clear intent to implement model cards, audit trails & lifecycle monitoring.
- AI access is restricted in customer-facing areas due to data-security policies, and users must complete training

Key Observations

What is going well?

- Clear **strategic intent** and visible **executive sponsorship**.
- Governance discipline emerging, with regulatory foresight (Consumer Duty, AI Act).
- Strong business demand and an active **pipeline of use cases**.
- Solid baseline of **tech investments** (Copilot, Instabase, CallMiner, Radar Live).

What needs focus?

- **Operationalizing AI as Business-as-Usual**: AI still sits in pilots and IT-led efforts; it is not embedded in day-to-day workflows.
- **Building Delivery Muscle**: Lack of small, empowered pods with product ownership, sprint cadence, and explicit RACI limits speed and impact.
- **Breaking Down Silos**: Initiatives remain fragmented across data, tools, and teams; decisions stall in committees and progress resets repeatedly.

Slide 85

AS1 I cannot fit the layout of other screens here
Anna Shymchenko, 2025-10-02T10:56:03.898

Current State Assessment

RSA – Current State SWOT Analysis

RSA's current-state SWOT analysis highlights a solid strategic foundation and emerging AI momentum, tempered by fragmented delivery, data integration challenges, and a critical need to translate early progress into scaled, differentiated outcomes.

Strength

- Strategic alignment under Intact Group, with clear focus on Commercial & Specialty growth.
- Early AI pilots in claims, underwriting, and productivity show willingness to experiment with new technology.
- Governance council and compliance guardrails in place, establishing a baseline for responsible AI adoption.
- Mature engineering practices (CI/CD, Microsoft ecosystem) provide a strong foundation for scaled digital delivery.

Weakness

- Underwriting & Ops: Manual re-keying across systems, inconsistent document handling & fragmented SOPs drive inefficiency & compliance risk.
- Claims: Reliance on free-text notes, manual triage/routing, slow FNOL coverage validation, missed hazard context, and fragmented supplier management.
- Sales & Distribution: Pipeline generation depends on manual trackers and portal scraping, with no automated propensity scoring.
- Data & Technology: Siloed datasets and poor integration block AI/ML adoption; early AI initiatives risk becoming one-off pilots without orchestration.
- People & Governance: Limited AI literacy, slow decision-making via large committees, and heavy HR reporting burdens reduce agility.

Opportunities

- Automation at scale: Streamline underwriting, renewals, sanctions checks, and claims triage with AI and workflow automation.
- Data enablement: Use governed micro-marts and short-term extracts to deliver value now, while converging onto Intact's enterprise model.
- Orchestration framework: Build standard connectors, guardrails, and observability to move beyond pilots and into enterprise AI adoption.
- Enhanced broker & customer engagement: Apply propensity scoring & real-time claims support to strengthen market differentiation.
- Risk prevention & recovery: Deploy AI for proactive monitoring of insured assets, fraud detection, and subrogation opportunities.
- People & culture uplift: Embed small empowered pods, improve AI literacy, and simplify governance processes to accelerate delivery.

Threat

- Margin erosion in Specialty if underwriting and claims processes remain slow and manual.
- Regulatory and compliance risks from inconsistent sanctions checks, policy coverage validation, and data quality.
- Operational risks from legacy platforms and fragmented systems causing errors, disputes, and audit gaps.
- Reputational risks from renewal delays, claims setup inefficiencies, and inconsistent broker experiences.
- Change fatigue and disengagement if AI remains stuck at pilot stage, while competitors advance with faster, more predictive solutions.

DP2

Current State Assessment

AI Technology

HEXAWARE

Hexaware's assessment indicates RSA's AI technology stack is in an early experimental phase — active at the application layer but constrained by missing orchestration, event-driven architecture, and compute infrastructure for scaled, agentic AI adoption.

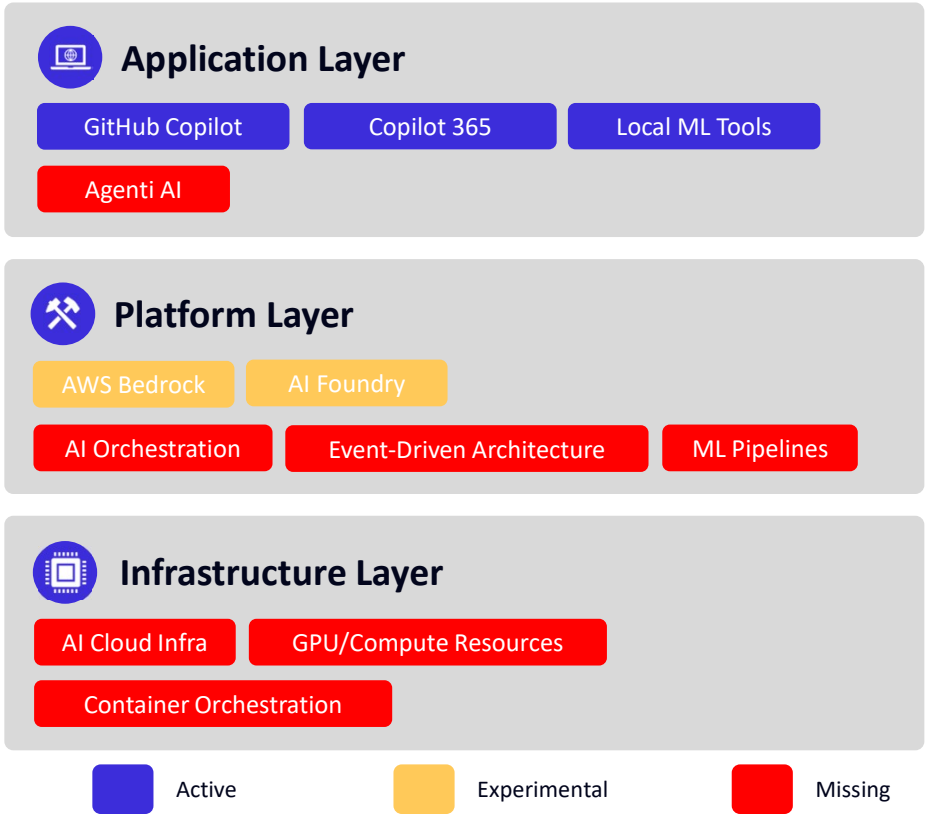

Architecture

Experimental phase with Intact's Labs materials using AWS Bedrock for foundational AI capabilities

Business

Operational ML experience with predictive models training using local Python tools based on company operational data

Individuals

Service-layer AI adoption through Microsoft ecosystem: GitHub Copilot, Copilot 365 for productivity enhancement

Current Gaps

- No infrastructure for agentic AI implementations
- No event-driven architecture
- Limited enterprise-scale AI enablement

Slide 87

AM1 [@Donna Pahel] I've removed the slide with your comment replacing it with this one that has our standard background

Anton Makarevich, 2025-09-26T08:47:20.405

DP2 [@Anton Makarevich] please review one-sentence intro here. Thanks.

Donna Pahel, 2025-10-08T17:35:55.563

AS2

Current State Assessment

HEXAWARE

Data Management & Integration

Hexaware's assessment indicates that RSA's data landscape remains highly fragmented, with siloed systems, inconsistent models, and limited governance impeding AI readiness and enterprise-wide insight generation. DP3

Data Silos

- Data is fragmented across multiple systems (PAS, claims systems, finance, broker management, e-trade workbenches, pricing engines, HR, etc.).
- Lack of centralized integration means each function operates with its own dataset, leading to duplication and inconsistency.

Manual & Inefficient Reporting

- HR, finance, and operational insights rely heavily on manual data pulls and offline spreadsheet manipulation.
- Limited automation in reporting pipelines → higher risk of errors and slow turnaround for insights.
- Analytics capabilities (BI, dashboards, predictive analytics) exist in pockets but are not standardized or enterprise-wide.

Lack of Data Integration

- Minimal integration between operational source systems and the central platform data ingestion is mostly point-to-point or manual (e.g., HR insights being done manually).

Data Governance Gaps

- Limited standardization of metadata, lineage, and quality controls.
- Business glossary and data dictionary exist in fragments, not widely adopted across teams.
- Ownership and stewardship of data are unclear in several domains.

Inconsistent Data Models

- No consistent enterprise-wide Enterprise Data Model (EDM) adopted.
- Project-specific or departmental logical/physical data models exist, but they are not integrated into a common structure.

Technology Landscape

- Legacy PAS and claims systems, broker tools, and actuarial/finance platforms coexist without seamless integration.
- Cloud adoption (e.g., Azure, Databricks) is visible in the target blueprint, but current usage is immature or siloed.

Slide 88

DP1 [@Dhanabal Subramaniam] and [@Ramprasad Jayaraman] – need POV content here
Donna Pahel, 2025-09-25T20:03:19.088

1 0 [@Donna Pahel] Modified the slide to include the details on the current state of data, based on the documents provided initially.
Dhanabal Subramaniam, 2025-09-29T13:00:48.497

AS2 Headline missing on this slide
Anna Shymchenko, 2025-10-02T11:21:24.060

DP3 [@Ramprasad Jayaraman] - please review the title and one-sentence intro here. Thanks.
Donna Pahel, 2025-10-08T17:38:00.745